Cluster-Adaptive Prompting for Multi-Turn Travel and Tourism Conversational Systems

Final Report

Surnella Sandeep

Department of Data Analytics, Walsh College

QM640: Data Analytics Capstone

Mr. Smitraj Raut

10th Nov 2026

GitHub repository: https://github.com/surnella/QM640-Dialogue-Clustering-Capstone

Contents

Abstract 4

Introduction 6

Background and Context 6

Problem Statement 7

Purpose of the Study and Research Questions 8

Contributions and Expected Value 9

Literature Review 9

Literature Review Approach and Relevance Matrix 9

Thematic Synthesis and Research Gaps 11

Materials and Methods 12

Data Sources, Access, and Ethical Governance 12

Data Dictionary and Data Cleaning 13

Exploratory Data Analysis (Univariate, Bivariate, and Multivariate) 15

Research Design, Hypotheses, and Statistical Power Analysis 18

Model Selection, Mathematical Formulation, and Evaluation Metrics 19

Architecture and Workflow 21

System Architecture and Algorithmic Workflows 21

Pipeline Component Specifications and Technology Stack 23

Results, Analysis, and Synthesis 23

Consolidated Model Performance and Hypothesis Testing Summary 23

RQ1 Results: Trajectory Clustering, Refined Archetypes, and Prefix Routing 25

RQ2 Results: Cluster-Adaptive Few-Shot Prompting on Primary Versus Refined Archetypes 30
RQ3 Results: 2 × 2 Factorial Evaluation of Declarative Rules and Adaptive Prompting .32
RQ4 Results: 2 × 3 Multi-Turn Simulation of Turn Economy Across Archetypes 35
Technical Interpretation, Business Significance, and Synthesis of Conclusions.....38
Implementation, Limitations, and Future Scope.....39
Deployment Architecture, System Integration, and Real-World Use Case.....39
Computational Reproducibility and Operational Requirements.....40
Study Limitations and Threats to Validity.....40
Future Scope and Research Extensions 41
References.....44
Appendix A.....49
Appendix B 52
Appendix C 53
Appendix D.....56
Appendix E 61
Appendix F.....65
Appendix G.....67

Abstract

Online travel agencies increasingly deploy large language model (LLM) assistants for vacation-package negotiation, yet one-size-fits-all prompts ignore customer decision-making heterogeneity while unguarded few-shot prompting risks entity hallucinations and policy violations. This study proposes and evaluates an integrated neuro-symbolic framework for archetype-aware, policy-governed, and turn-efficient conversational artificial intelligence (AI) in travel booking. Evaluated on 1,334 Microsoft Frames dialogues (70% golden discovery, $n = 933$; 30% held-out test, $n = 401$) and 1,499 MultiWOZ 2.2 comparison dialogues, the research follows a four-stage progression: **Identify (RQ1) → Personalize (RQ2) → Govern (RQ3) → Optimize (RQ4)**. First, to **identify** customer interaction styles early (**RQ1**), clustering algorithms differed significantly across 159 runs (53 matched $n = 500$ subsamples), one-way $F(2, 156) = 91.33$, $p < .001$, $\eta^2 = .54$ (matched repeated-measures $F(1.72, 89.30) = 83.04$, Friedman $\chi^2(2) = 56.30$, $p < .001$). Selecting *K*-Means ($k = 3$) for 100% routing coverage (versus HDBSCAN's 39.1% MiniLM / 35.8% Gemini noise rejection) on a 46-dimensional hybrid *Blocks* representation over the first nine user messages established three stable reference archetypes (bootstrap ARI = .89 in the truncation grid; .87 in the 6-configuration ablation): *Decisive Booker*, *Open Explorer*, and *Rigid-Constraint Seeker*. Prefix Logistic Regression routers achieved Cohen's $\kappa = .535$ (70.1% agreement) after five user messages and $\kappa = .958$ (97.3% same-feature consistency) after nine messages against the nine-message reference labels. Second, to **personalize** assistant communication (**RQ2**, $N = 33$ repeated-measures dialogues), routing via the five-message *Blocks* router into a combined package of archetype-partitioned FAISS sub-indices and pacing directives improved de-lexicalized SBERT style similarity from .453 (static few-shot) and .530 (zero-shot) to .597, $F(1.88, 60.29) = 9.79$, $p < .001$, $\eta_p^2 = .23$, and cut pacing error from 10.12

words (zero-shot) and 6.33 words (static) to 3.87 words, $F = 7.10$, $p = .007$, whereas generic semantic metrics stayed invariant and coarse whole-dialogue routing ($\kappa = .24$) had yielded null effects ($p = .211$). Third, to **govern** business and factual correctness (**RQ3**, 2 × 2 factorial, $N = 144$ evaluations across 36 dialogues; dual-judge $\kappa = .917$), pairing cluster-adaptive prompting with a declarative rule engine, [SLOT] masking, and leakage sanitization raised violation enforcement sensitivity from 5.6% to 97.2% ($F(1, 68) = 498.78$, $p < .001$; compliant specificity = 88.9% vs. 94.4%, $p = .408$; overall rule accuracy = .931, $F(1, 140) = 43.92$, $p < .001$), boosted Slot F_1 to .952 ($F = 489.52$, $p < .001$), and eliminated the .359 (55.6% incidence) unguarded exemplar leakage rate to .000 (interaction $F(1, 140) = 34.82$, $p < .001$; Welch $t(35) = 5.90$, $p < .001$) while retaining .506 style similarity. Fourth, to **optimize** multi-turn booking efficiency (**RQ4**, 2 × 3 simulation, $N = 180$ unique seed dialogues, $n = 30$/cell; persona $\kappa = .933$), a *Verified Informed Slot Ledger* and archetype-specific policies reduced total turns (user plus assistant messages) to package agreement by 42.7% for *Rigid-Constraint Seekers* (8.73 → 5.00, reaching the 5-turn protocol floor), 34.2% for *Decisive Bookers* (7.60 → 5.00), and 29.5% for *Open Explorers* (7.47 → 5.27; $F(1, 174) = 664.43$, interaction $F(2, 174) = 10.89$, both $p < .001$) while eliminating redundant slot re-asking (.109 → .000). Combining behavioral routing, adaptive communication, explicit policy controls, and state-aware dialogue management thus supports personalized, compliant, and efficient travel-booking interactions.

Keywords: conversational artificial intelligence, customer archetypes, large language models, few-shot prompting, business-rule enforcement, slot filling, dialogue management, travel booking

Introduction

Background and Context

Conversational artificial intelligence (AI) assistants have become a core digital channel in travel, tourism, and hospitality, supporting destination discovery, package comparison, and itinerary booking (Buhalis & Sinarta, 2019; Ivanov & Webster, 2019; Jannach et al., 2021). Unlike single-turn question answering, vacation package booking is a consultative, multi-turn negotiation in which travelers balance competing constraints across origin, destination, dates, duration, hotel tier, and budget (El Asri et al., 2017). Empirical analyses of human travel dialogues demonstrate that travelers converse in distinct behavioral styles: some state fixed parameters up front and seek immediate confirmation, others explore multiple destinations and hotel tiers side by side, and others insist on rigid date and luxury constraints that require proactive trade-off resolution when preferred inventory exceeds their budget ceiling (Inaba et al., 2022; Melián-González et al., 2021).

Despite this behavioral diversity, commercial large language model (LLM) travel assistants typically use one static system prompt and a uniform sequential intake script. Asking clarifying questions one slot at a time or offering a single hotel option per turn creates repetitive turns and session abandonment. Conversely, adapting an LLM's style purely through in-context few-shot examples introduces two severe reliability hazards (Liu et al., 2022; Rebedea et al., 2023): first, unmasked retrieved examples cause factual entity leakage, where the model copies stale hotel names, airfares, or promo codes from the examples into the active session; second, stylistic prompt instructions compete with business rules, causing the assistant to overlook mandatory visa windows, passport validity thresholds, party-size limits, or discount caps (Yao et

al., 2024). Resolving both personalization and compliance requires a neuro-symbolic architecture that decouples behavioral style adaptation from deterministic policy verification.

Problem Statement

Business problem statement. Online travel agencies (OTAs), tour operators, and hospitality contact centers suffer from low self-service conversion and elevated session abandonment because static conversational assistants fail to adapt their information density and recommendation strategy to distinct customer decision-making styles, while unverified generative responses expose enterprises to financial and regulatory liability from booking policy violations and hallucinated package terms (Buhalis & Sinarta, 2019; Ivanov & Webster, 2019; Melián-González et al., 2021).

Technical problem statement. Existing LLM dialogue systems lack (a) an empirically validated representation for clustering multi-turn user behavior into stable archetypes that can be routed accurately from early dialogue prefixes without end-of-dialogue outcome leakage, (b) a slot-masked exemplar retrieval mechanism that achieves target archetype style convergence without inducing entity leakage, and (c) formal factorial evidence quantifying the main and interaction effects of external declarative rule verification and cluster-adaptive prompting on slot fidelity, policy adherence, and multi-turn turn economy (Acikgoz et al., 2025; Liu et al., 2022; Rebedea et al., 2023).

Measurable research formulation. This study models six quantitative dependent variables (Y)—cluster separation Silhouette (Y_1), style convergence cosine and pacing error (Y_2), policy rule adherence rate (Y_3), exemplar entity leakage rate (Y_4), factual slot F_1 (Y_5), and log-transformed turns to goal completion and redundant elicitation rate (Y_6)—as functions of four manipulated independent variables (X): clustering algorithm and trajectory representation (X_1,

RQ1), in-context conditioning strategy (X_2, RQ2), rule-engine mode crossed with prompting strategy (X_3, RQ3), and prompting strategy crossed with refined user archetype (X_4, RQ4).

Purpose of the Study and Research Questions

The purpose of this quantitative study is to develop and evaluate an end-to-end neuro-symbolic, cluster-adaptive conversational framework on public human travel dialogues across four linked research questions (Table 1).

Table 1

Research Questions, Hypotheses, and Analytical Framework

RQ	Research question	Null and alternative hypotheses	Experimental design and statistical test
RQ1	Which clustering algorithm (K-Means, Ward agglomerative, HDBSCAN) best separates dialogue embeddings, and which stable, early-routable user archetypes emerge?	H_{01}: Equal mean Silhouette across algorithms. H_{a1}: At least one differs, and truncated behavioral-semantic prefixes yield stable (ARI ≥ .60), routable archetypes.	Resampling benchmark across N = 159 runs (53 matched n = 500 subsamples: one-way ANOVA, matched RM-ANOVA, Friedman, Welch, Kruskal–Wallis, Tukey & Games–Howell) + truncation grid
RQ2	Does cluster-adaptive few-shot prompting increase style convergence and pacing calibration relative to zero-shot and static few-shot prompting?	H_{02}: Equal mean style cosine across conditions. H_{a2}: Cluster-adaptive > static and zero-shot on refined user archetypes.	One-way repeated-measures ANOVA (N = 33 dialogues × 3 conditions = 99 calls; Greenhouse–Geisser, Friedman, Holm contrasts)
RQ3	How do a declarative rule engine and cluster-adaptive prompting interact to govern rule adherence, exemplar leakage, slot F_1, and style alignment?	H_{03}: No main or interaction effects. H_{a3}: Active rules raise adherence and slot F_1 and eliminate few-shot entity leakage.	2 × 2 factorial ANOVA (N = 144 across 36 held-out dialogues, n = 36/cell; OLS, HC3, ART, sensitivity/specificity split; dual-LLM judge κ)
RQ4	Does cluster-adaptive prompting with an informed-slot ledger reduce turns to goal completion and redundant elicitation differently across user archetypes?	H_{04}: No prompting or prompting × archetype interaction effects. H_{a4}: Cluster-adaptive prompting reduces log turns and RER with an archetype interaction.	2 × 3 between-subjects ANOVA on log turns and RER (N = 180 unique held-out dialogues, n = 30/cell; OLS, HC3, ART, Welch t, Mann–Whitney U; persona κ)

Note. RQ = research question; ARI = adjusted Rand index; HC3 = heteroscedasticity-consistent SE; ART = aligned rank transform; RER = redundant elicitation rate. All tests use α = .05.

Contributions and Expected Value

This study delivers four contributions: (1) demonstrating that truncating dialogues to the first nine user messages and combining Gemini prefix embeddings with structured LLM behavioral ratings (*Blocks*) yields three stable reference archetypes (bootstrap ARI = .89 in the truncation grid; .87 in the 6-configuration ablation) with chance-corrected held-out prefix routing agreement of $\kappa = .535$ (70.1% raw agreement) at five user messages and $\kappa = .958$ (97.3% same-feature boundary consistency) at nine user messages against the unsupervised nine-message *K*-Means reference clusters; (2) contrasting primary versus refined archetypes in RQ2 to demonstrate that combining archetype-partitioned exemplar retrieval with routed archetype pacing directives significantly improves de-lexicalized SBERT structural style similarity ($p < .001$) and archetype pacing calibration ($p = .007$) once upstream routing is coherent, while generic semantic similarity remains unchanged; (3) demonstrating in RQ3 that an integrated governance layer (declarative rule engine, [SLOT] masking, and post-generation sanitization) eliminates the 35.9% entity leakage (55.6% incidence) of unguarded few-shot prompting in the evaluated sample while raising policy violation enforcement sensitivity from 5.6% to 98.6% and preserving 88.9% compliant-case acceptance specificity; and (4) showing across 180 independent multi-turn simulations in RQ4 that coupling the *Verified Informed Slot Ledger* with archetype-specific turn-economy policies reduces total conversational turns to booking agreement by 29.5% to 42.7% and eliminates redundant slot questions.

Literature Review

Literature Review Approach and Relevance Matrix

Peer-reviewed studies were retrieved from the ACL Anthology, ACM Digital Library, IEEE Xplore, ScienceDirect, Google Scholar, and arXiv using Boolean combinations of *travel*

and hospitality chatbots, *conversational recommender systems*, *task-oriented dialogue corpora*, *trajectory clustering*, *in-context exemplar retrieval*, and *LLM guardrails*. Twenty-six sources spanning business adoption, dialogue datasets, unsupervised clustering, retrieval-augmented prompting, and statistical methodology were synthesized. Table 2 maps the core literature streams to the gaps addressed in RQ1–RQ4.

Table 2

Literature Relevance Matrix Across Business and Technical Domains

Study and stream	Methodology and dataset	Key findings	Gap addressed in this study
Buhalis & Sinarta (2019); Ivanov & Webster (2019); Melián-González et al. (2021) — Travel AI business value	Service-automation frameworks and consumer adoption modeling (N = 1,476) in tourism and hospitality	Real-time personalization and low interaction effort drive chatbot booking conversion and ROI	Measures perceptions via surveys rather than instrumenting turn economy, slot redundancy, or policy compliance in LLM dialogues (RQ3, RQ4)
Jannach et al. (2021); Inaba et al. (2022) — Conversational recommenders & traveler styles	CRS architectural taxonomy; collection and analysis of multi-turn travel agency dialogues	Travelers differ in initiative, exploration breadth, and constraint rigidity; static policies cause repetitive turns	Describes traveler diversity qualitatively without unsupervised archetype clustering or prefix routing (RQ1, RQ2)
El Asri et al. (2017); Budzianowski et al. (2018); Zang et al. (2020) — Dialogue corpora	Microsoft Frames (1,369 package negotiations) and MultiWOZ 2.2 (10,438 multi-domain dialogues)	Frames captures nonlinear frame-switching and budget trade-offs; MultiWOZ standardizes slot tracking	Built for supervised state tracking; neither clusters user archetypes nor evaluates few-shot leakage or policy rules (RQ1–RQ4)
Reimers & Gurevych (2019); Lee et al. (2025); McInnes et al. (2017, 2018); Hennig (2007) — Embeddings & clustering	Sentence-BERT (384-d), Gemini embeddings (3,072-d), HDBSCAN, UMAP, and bootstrap ARI stability	Cosine geometry supports semantic clustering; bootstrap ARI and Silhouette quantify partition stability	Clusters static texts rather than prefix-truncated user trajectories paired with LLM behavioral ratings (RQ1)
Brown et al. (2020); Liu et al. (2022); Johnson et al. (2021) — In-context learning & FAISS	Few-shot prompting; k-NN exemplar selection (KATE); FAISS vector indexing	Nearest-neighbor in-context examples outperform random static examples on single-turn NLP tasks	Retrieves from one global pool without archetype partitioning, slot masking, or style/leakage evaluation (RQ2, RQ3)

Study and stream	Methodology and dataset	Key findings	Gap addressed in this study
Rebedea et al. (2023); Yao et al. (2024); Acikgoz et al. (2025); Gilardi et al. (2023) — Guardrails & evaluation	NeMo Guardrails; τ-bench policy benchmark; TD-EVAL metrics; LLM annotation reliability	Prompt-only LLMs violate domain rules; programmable rails and turn-level factual audits are essential	Does not test the 2 × 2 factorial interaction between declarative rules and cluster-adaptive style prompting (RQ3)

Note. CRS = conversational recommender system; ARI = adjusted Rand index; FAISS = Facebook AI Similarity Search.

Thematic Synthesis and Research Gaps

Business context: Hospitality service automation and conversational friction. In digital travel and hospitality commerce, competitive differentiation depends on real-time itinerary co-creation ('nowness' service; Buhalis & Sinarta, 2019). Ivanov and Webster (2019) showed that automated service agents generate financial value only when they resolve consultative inquiries without human escalation, while Melián-González et al. (2021) demonstrated that consumer intention to use travel chatbots depends on interaction efficiency and task fit. Empirical studies of travel agency dialogues confirm that human consultants adapt their pacing—confirming bookings rapidly for decisive buyers while presenting side-by-side comparisons for exploratory shoppers (Inaba et al., 2022; Jannach et al., 2021).

Task-oriented corpora and unsupervised trajectory clustering. While MultiWOZ (Budzianowski et al., 2018) and MultiWOZ 2.2 (Zang et al., 2020) benchmark slot tracking across local services, they lack flight routing and multi-package price negotiation. Microsoft Frames (El Asri et al., 2017) directly captures multi-turn flight-and-hotel negotiation under budget constraints, making it the ideal primary corpus and MultiWOZ 2.2 an external structural comparison. Representing dialogues with Sentence-BERT (Reimers & Gurevych, 2019) and Gemini embeddings (Lee et al., 2025) enables rigorous comparison of K-Means, Ward agglomerative clustering, and HDBSCAN (McInnes et al., 2017) using Silhouette (Rousseeuw,

1987), Davies–Bouldin (Davies & Bouldin, 1979), bootstrap ARI stability (Hennig, 2007; Hubert & Arabie, 1985), and UMAP (McInnes et al., 2018; Pedregosa et al., 2011), augmented by reliable zero-temperature LLM behavioral ratings (Gilardi et al., 2023).

Dynamic in-context learning, neuro-symbolic guardrails, and research gaps. Few-shot prompting (Brown et al., 2020) and similarity-based FAISS retrieval (Johnson et al., 2021; Liu et al., 2022) improve single-turn task accuracy but risk entity contamination and policy drift in multi-turn dialogues. Programmable guardrails (Rebedea et al., 2023), policy benchmarks (τ-bench; Yao et al., 2024), and turn-level evaluation frameworks (TD-EVAL and BERTScore; Acikgoz et al., 2025; Zhang et al., 2020) motivate decoupling style from policy verification. Four research gaps follow: (1) discovering early-routable travel user archetypes free of outcome confounding (RQ1); (2) testing archetype-matched, slot-masked exemplar retrieval for style convergence (RQ2); (3) isolating the 2 × 2 factorial interaction between declarative rule verification and adaptive prompting (RQ3); and (4) quantifying multi-turn turn economy and redundant elicitation reductions across traveler personas (RQ4).

Materials and Methods

Data Sources, Access, and Ethical Governance

This study uses two public English dialogue corpora (Table 3): Microsoft Frames (El Asri et al., 2017; 1,369 raw dialogues, 1,334 analytic) and MultiWOZ 2.2 (Zang et al., 2020; seeded sample of 1,500 travel dialogues, 1,499 analytic). Both datasets were collected with research consent, use anonymized role-player IDs, and contain no personally identifiable information.

Table 3

Dataset Overview

Attribute	Microsoft Frames (primary corpus)	MultiWOZ 2.2 (comparison corpus)
Records	1,369 raw; 1,334 analytic (golden D_{gold} = 933, held-out D_{test} = 401)	10,438 raw; travel sample of 1,500; 1,499 analytic
Domain setting	Multi-turn flight + hotel package negotiation under budget and date constraints	Multi-domain city services (hotel, train, taxi, attraction) without flights
Annotations used	Utterance text, speaker, dialogue acts, slot–value arguments, wizard task success, user rating	Utterance text, speaker, active services, turn-level slot state dictionaries
Analytical role	RQ1 clustering & refined archetypes, RQ2 prompting, RQ3 2×2 experiment, RQ4 2×3 simulation	External descriptive comparison in EDA and cross-corpus projection check in RQ1

Data Dictionary and Data Cleaning

Both corpora were normalized into a two-tier Pydantic schema (*src/canonical_schema.py*): Tier 1 stores nine universal travel slots (origin, destination, departure/return dates, duration in nights, adults, children, budget in USD, and hotel tier), and Tier 2 stores five policy slots evaluated by the rule engine (visa window in days, passport validity in months, package type, room sharing, and discount percentage). Table 4 defines the analytical variables and units (raw source fields and canonical schema fields are listed in Appendix A, Table A3). Table 5 documents the programmatic quality audit and cleaning log. To prevent data leakage, the 1,334 analytic Frames dialogues were split once (seed 42, stratified on task success) into a 70% golden discovery partition (D_{gold}, n = 933) used to fit clusters, PCA transformations, routers, and FAISS indices, and a 30% held-out test partition (D_{test}, n = 401) used strictly for out-of-sample evaluation.

Table 4

Data Dictionary of Analytical Variables With Units and Operational Definitions

Variable	Role	Type and unit	Operational definition and scale	RQ
Clustering algorithm	Primary IV	Categorical (3 levels)	K-Means (k=3), Ward agglomerative (k=3), HDBSCAN	1
Silhouette & Davies–Bouldin	Primary/Sec. DV	Continuous (unitless / ratio)	Cosine Silhouette [−1, 1] and cluster scatter ratio (≥ 0)	1
Adjusted Rand index (ARI)	Diagnostic DV	Continuous [−1, 1] (unitless)	Chance-corrected partition stability across 80% subsamples	1
LLM behavioral profile (16 fields)	Clustering / Audit	10 ordinal (1–5) + 6 cat./bool.	10 Likert scales + 3 goal/tone clustering fields + 3 audit fields (Table C1)	1, 4
Refined user archetype	IV / Moderator	Categorical (3 levels)	*Decisive Booker*, *Open Explorer*, *Rigid-Constraint Seeker*	1–4
Prompting strategy	Primary IV	Categorical (2 or 3 levels)	Zero-Shot, Static Few-Shot (k=3), Cluster-Adaptive Few-Shot (k=3)	2–4
Style convergence & Pacing err.	Primary/Sec. DV	Cosine [−1, 1] / Words (count)	Cosine to target archetype agent centroid; $\lvert words_{gen} - \mu_{words,a} \rvert$	2, 3
Rule-engine mode	Primary IV	Binary (2 levels)	Rules Off (unguarded) vs. Rules Active (pre-check + verification)	3
Rule adherence rate (RAR)	Primary DV	Proportion [0, 1] (rate)	Share of violated policy rules enforced (and compliant states unblocked)	3
Exemplar leakage rate (ELR)	Primary DV	Proportion [0, 1] (rate)	Share of response entities originating from prompt exemplars	3
Factual slot coverage F_1	Primary DV	Proportion [0, 1] (F_1 score)	Harmonic mean of precision and recall over active booking state slots	3
Turns to goal (log) & RER	Primary/Sec. DV	Log count (ln turns) / Rate [0, 1]	Log turns until [ACCEPT_BOOKING]; redundant slot requests / agent turns	4

Note. IV = independent variable; DV = dependent variable; RER = redundant elicitation rate.

Raw and canonical fields are in Table A3.

Table 5

Data Cleaning and Quality Audit Log

Issue	Variables	Detection	Frames	MultiWOZ	Treatment and rationale
Duplicate full text	Utterances	SHA-256 text hash	0	1	Removed duplicate to prevent double-weighting
Fewer than 3 turns	total_turns	Turn count < 3	35	0	Excluded; minimum 3 turns needed for prefix routing
Turn-count outliers	total_turns	> Q3 + 1.5 × IQR	22 (> 32.5)	13 (> 27.0)	Retained as genuine negotiations; log-transformed in RQ4

Issue	Variables	Detection	Frames	MultiWOZ	Treatment and rationale
Unobserved slots	Tier 1 & 2 slots	Never mentioned	Table E1	Table E1	Preserved as null (None); never imputed
Entity values in text	Utterance text	Slot/number regex	All	All	Masked with [SLOT] tokens in exemplars to prevent leakage
Analytic sample	All fields	Post-cleaning count	1,334 (933/401)	1,499	Locked with stratified 70/30 split (seed 42)

Note. Zero duplicate IDs or empty utterances occurred in either corpus. Missing user ratings (11 in Frames) were retained.

Exploratory Data Analysis (Univariate, Bivariate, and Multivariate)

Univariate Analysis

Table 6 and Figure 1 summarize the univariate distributions. Frames and MultiWOZ 2.2 have similar mean lengths (M = 14.93 vs. 15.17 turns), but Frames exhibits greater dispersion (SD = 6.74 vs. 4.46) and a heavier right tail due to multi-offer package comparisons. Golden and held-out partitions are nearly identical across all metrics. Slot completeness (Appendix E, Table E1 and Figure E1) confirms high coverage of origin (98.4%), destination (99.1%), dates (96.0%), and budget (78.5%) in Frames, but 0.0% coverage of visa, passport, and discount rules in both corpora, motivating synthetic Tier-2 injection in RQ3.

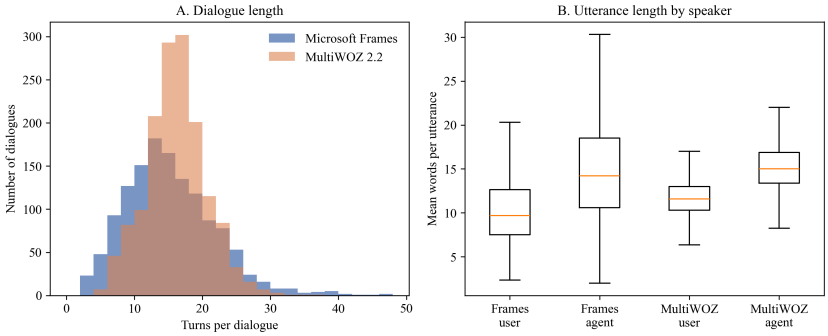
Table 6

Univariate Descriptive Statistics of the Analytic Samples

Sample	*N*	Turns *M* (*SD*)	*Mdn*	Range	User words	Agent words	Task success (%)
Microsoft Frames	1,334	14.93 (6.74)	14	3–48	10.6	15.0	94.2
MultiWOZ 2.2	1,499	15.17 (4.46)	16	4–36	11.7	15.1	—
Frames golden partition (D_gold)	933	14.95 (6.75)	14	3–48	10.7	15.3	94.2
Frames held-out partition (D_test)	401	14.90 (6.71)	14	3–43	10.4	14.5	94.3

Figure 1

Univariate Distribution of Dialogue Length by Corpus



Note. Frames $N = 1{,}334$; MultiWOZ $N = 1{,}499$.

Bivariate Analysis

Parametric (Pearson r) and non-parametric (Spearman ρ) bivariate correlations across all 1,334 Frames dialogues (Table 7; Figure 2) show that dialogue length is strongly driven by slot revisions ($r = .74$, $\rho = .76$, $p < .001$) and agent package offers ($r = .56$, $\rho = .61$, $p < .001$). Linking structural counts to LLM behavioral ratings confirms construct validity: slot revisions correlate strongly with user exploration breadth ($r = .60$, $\rho = .67$, $p < .001$), whereas decisiveness correlates negatively with dialogue length ($r = -.31$, $\rho = -.29$, $p < .001$).

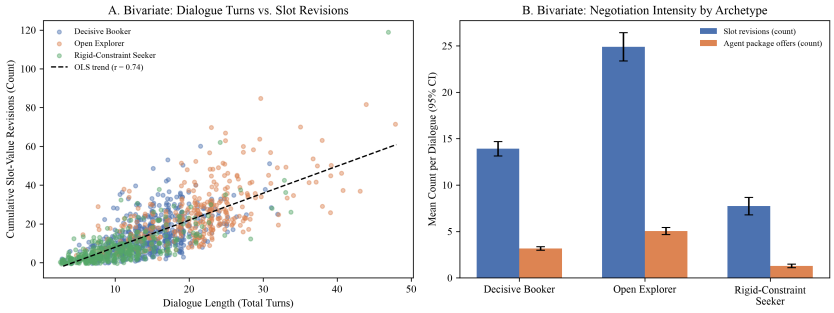
Table 7

Bivariate Parametric and Non-Parametric Correlations of Dialogue Dynamics (Frames)

Variable pair	Pearson r	p (Pearson)	Spearman ρ	p (Spearman)
Total Turns vs. Slot Revisions	.741	< .001	.759	< .001
Total Turns vs. Agent Package Offers	.561	< .001	.608	< .001
Total Turns vs. User Hesitation Markers	.336	< .001	.355	< .001
Slot Revisions vs. Exploration Breadth Rating	.603	< .001	.673	< .001
Total Turns vs. User Decisiveness Rating	−.310	< .001	−.289	< .001
User Words/Turn vs. Information Up Front Rating	.158	< .001	.206	< .001
Agent Offers vs. Destination Flexibility Rating	.324	< .001	.358	< .001

Figure 2

Bivariate Relationships Between Dialogue Length, Slot Revisions, and Archetype Negotiation Intensity



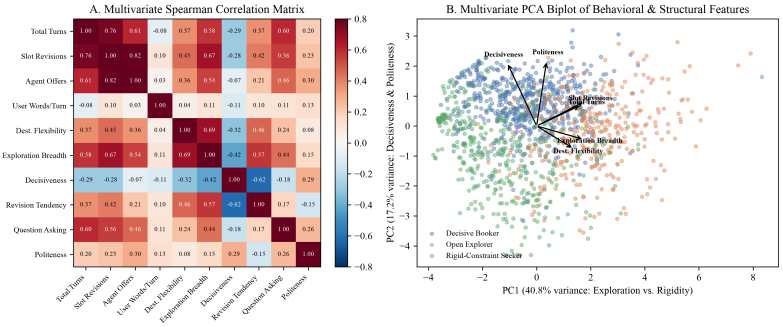
Note. Panel A: Turns vs. slot revisions with OLS trend. Panel B: Mean revisions and package offers by refined archetype (95% CI).

Multivariate Analysis

Multivariate structure across ten structural and behavioral variables was examined via Spearman rank correlation and principal component analysis (PCA; Figure 3; full loadings in Appendix E, Table E2). The first three components explain 68.6% of total variance: PC1 (40.8%) captures *Exploration vs. Rigidity* (revisions .42, turns .39, exploration breadth .39), PC2 (17.2%) captures *Decisiveness & Politeness* (−.51, −.46), and PC3 (10.6%) isolates *User Verbosity* (.89), separating the three refined archetypes across a continuous behavioral manifold.

Figure 3

Multivariate Spearman Correlation Matrix and Principal Component Biplot of Behavioral Features



Note. Panel A: Rank correlation matrix. Panel B: PCA biplot colored by refined archetype with top loading vectors.

Research Design, Hypotheses, and Statistical Power Analysis

Sample sizes were computed a priori for each planned analysis using Cohen's medium effect $f = .25$ (Cohen, 1988), $\alpha = .05$, and target power $1 - \beta \geq .80$ via the noncentral F distribution (Faul et al., 2007):

$$1 - \beta = 1 - F_{nc}(F_{crit};\, df_1,\, df_2,\, \lambda), \quad F_{crit} = F_{1-\alpha}(df_1,\, df_2) \qquad (1)$$

where $\lambda = f^2 N$ for between-subjects designs (RQ1, RQ3, RQ4) and $\lambda = f^2 N m \varepsilon / (1 - \rho)$ for repeated-measures RQ2 ($m = 3$, $\rho = .50$, $\varepsilon = 1$). Minimum sample sizes were rounded to balanced cells with a 10% buffer (Table 8), yielding final power of .805–.878. In RQ1, because each of the 53 random seeds draws $n = 500$ dialogues without replacement from D_{gold} ($n = 933$, a 53.6% finite-population fraction), we pair the 159-run one-way benchmark with a matched-subsample repeated-measures ANOVA and Friedman test across the 53 shared seeds plus full-D_{gold} point estimates. In RQ3, 36 stratified held-out dialogues (18 policy-violating, 18 compliant; 12 per archetype) are evaluated across all four 2 × 2 cells ($N = 144$ evaluations). In RQ4, **180 distinct, non-overlapping held-out seed dialogues** (60 per archetype, partitioned into $n = 30$ unique

dialogues for *Zero-Shot* and $n = 30$ unique dialogues for *Cluster-Adaptive Few-Shot*) are allocated across the six 2 × 3 cells, ensuring strict between-subjects independence ($df_{error} = 174$).

Table 8

Sample Size and Power Calculations Matched to Each Planned Analysis

RQ	Planned analysis	*df* at N_{min}	λ	F_{crit}	N_{min}	Power	N_{final} (per group)	Final power
RQ1	One-way ANOVA (3 algorithms)	2, 155	9.88	3.05	158	.802	159 (53)	.805
RQ2	One-way repeated-measures ANOVA (3 conditions)	2, 54	10.50	3.17	28	.812	33 (11)	.878
RQ3	2 × 2 factorial ANOVA	1, 124	8.00	3.92	128	.801	144 (36)	.846
RQ4	2 × 3 factorial ANOVA	2, 152	9.88	3.06	158	.802	180 (30)	.855

Model Selection, Mathematical Formulation, and Evaluation Metrics

Trajectory and hybrid user representations (RQ1). Primary trajectory vectors pool de-lexicalized turn embeddings $\mathbf{e}_t$ with exponential decay ($\lambda_d = 0.15$):

$$\mathbf{v} = \Sigma_{t=1}^{T} w_t \mathbf{e}_t, \quad w_t = \exp[-\lambda_d(T - t)] \,/\, \Sigma_{j=1}^{T} \exp[-\lambda_d(T - j)] \quad (2)$$

To prevent end-of-dialogue outcome confounding, the refined archetype representation (*Blocks*) concatenates two L2-normalized blocks computed strictly on the first M user messages: (1) the top 20 principal components (fitted on D_{gold}) of the 3,072-dimensional Gemini embedding (*gemini-embedding-001*; Lee et al., 2025), and (2) a standardized 16-field LLM behavioral schema comprising 13 clustering fields (ten 1–5 ordinal rating scales plus primary goal, secondary goals, and emotional tone, encoded into 26 numeric columns) and three audit-only fields (session outcome, party type, and unusual-conversation flag, excluded from clustering input; Table C1; test–retest weighted κ = .87–.96 across the 10 ordinal scales and κ = .75–.98 across the 6 categorical/boolean fields; revision correlation ρ = .73, Table C2). In RQ2–RQ4, *RefinedArchetypeRouter* projects each test dialogue's opening user prefix into this golden-fitted 46-dimensional *Blocks* space, predicts the user archetype via L1-regularized logistic regression,

and applies a combined adaptation package (retrieving $k = 3$ slot-masked exemplars from that archetype's FAISS sub-index alongside an explicit archetype pacing/action directive). In RQ3, the deterministic rule engine (*src/rule_engine.py*, 14/14 boundary tests passed) evaluates five YAML policy rules (Table 9) within an integrated governance pipeline combining pre-generation rule injection, [SLOT] exemplar masking, and post-generation regex sanitization. In RQ4, each closed-loop simulation pairs the *Verified Informed Slot Ledger* with archetype-specific turn-economy policies, counting total conversational turns (user plus assistant messages) until the simulated traveler issues <code>[ACCEPT_BOOKING]</code> (subject to a protocol minimum of 5 total turns: <code>U1 → A1 → U2 → A2 → U3</code> once budget is confirmed). Table 10 defines all evaluation metrics.

Table 9

Declarative Policy Rules Used in RQ3

Rule ID and name	Trigger condition	Action	Severity
RULE_VISA_WINDOW: Visa Processing Lead Time	departure_lead_days < 15; is_international = true	block and flag	fatal
RULE_PASSPORT_VALIDITY: Passport Validity Threshold	passport_expiry_months < 6; is_international = true	block and flag	fatal
RULE_PAX_CAPACITY: Maximum Adult Capacity per Single Booking	traveler_count_adults > 8; package_type = 'FIT'	convert to group	warning
RULE_DISCOUNT_THRESHOLD: Maximum Discretionary Sales Discount	discount_applied_percentage > 15.0	clamp discount	fatal
RULE_BUDGET_MINIMUM: Minimum Feasible Per-Pax International Budget	budget stated; budget_amount < 500; is_international = true	prompt realistic budget	warning

Table 10

Evaluation Metrics, Mathematical Formulae, and Interpretation

Metric	Formula	RQ	Interpretation
Silhouette & DB index	$s(i) = [b(i)-a(i)]/\max\{a,b\}$; $DB = (1/k)\Sigma \max(\sigma_i+\sigma_j)/d(\mathbf{c}_i,\mathbf{c}_j)$	1	Higher Silhouette [−1, 1] and lower DB (≥ 0) indicate better separation
ARI & Cohen's κ	$ARI = (RI-\mathrm{E}[RI])/(\max RI-\mathrm{E}[RI])$; $\kappa = (p_o-p_e)/(1-p_e)$	1, 3, 4	Bootstrap 80%-subsample stability (ARI ≥ .60) and routing/judge agreement
Style cosine & Pacing err.	$\cos(\mathbf{r}, \mathbf{c}_a) = (\mathbf{r}\cdot\mathbf{c}_a)/(\|\mathbf{r}\|\|\mathbf{c}_a\|)$; $\lvert W_{\text{gen}} - \mu_{W,a}\rvert$	2, 3	De-lexicalized SBERT structural style cosine and word-count pacing error
Rule adherence (RAR)	Sens. = $\lvert R_{\text{enf}} \cap R_{\text{viol}}\rvert/\lvert R_{\text{viol}}\rvert$ (n=18/cell); Spec. = 1 − false_block (n=18/cell); Overall (n=36/cell)	3	Separated into Violation Enforcement (Sensitivity), Compliant Acceptance (Specificity), and Overall Accuracy
Exemplar leakage (ELR)	$\lvert E_{\text{resp}} \cap (E_{\text{ex}} \setminus E_{\text{diag}})\rvert$ / $\max(1, \lvert E_{\text{resp}}\rvert)$	3	Proportion [0, 1] of response entities leaked from prompt exemplars
Slot coverage F_1	2 · Precision · Recall / (Precision + Recall) over active state slots	3	Proportion [0, 1] of active booking slots grounded in response
Turns to goal & RER	TGC = total turns (user + assistant messages) to [ACCEPT_BOOKING] (min = 5); RER = $A_{\text{redundant}}$ / A_{total}	4	Lower total/log conversational turns and lower RER indicate superior turn economy

Note. Parametric ANOVAs are paired with Welch/HC3 and non-parametric (Kruskal–Wallis, Friedman, Mann–Whitney, ART; Wobbrock et al., 2011) checks.

Architecture and Workflow

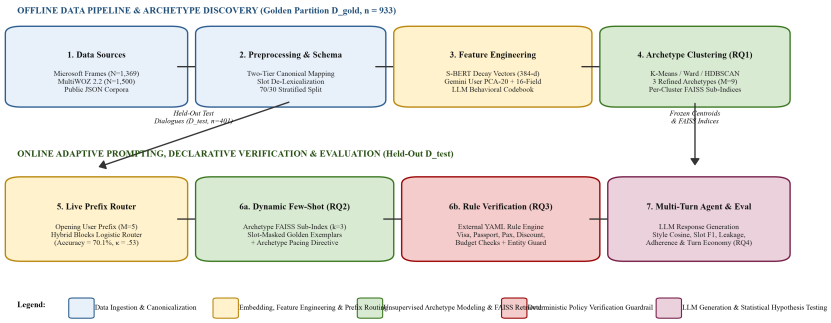
System Architecture and Algorithmic Workflows

Figure 4 illustrates the five-layer neuro-symbolic system architecture, and Figure 5 details the step-by-step execution workflows for RQ1 through RQ4. Offline discovery and indexing (Layers 1–2) operate strictly on the golden partition (D_{gold}, n = 933): raw dialogues are parsed into two-tier Pydantic trajectories, prefix embeddings (*gemini-embedding-001*) and 16-field LLM behavioral profiles are combined into the *Blocks* representation on the first M = 9 user messages, and three K-Means centroids, three archetype-specific FAISS *IndexFlatIP* sub-indices,

and the L1-regularized prefix router are frozen. Online runtime and evaluation (Layers 3–5) operate on held-out test dialogues (D_{test}, $n = 401$): incoming prefixes are projected into the golden *Blocks* space and routed to their predicted archetype, $k = 3$ slot-masked exemplars are retrieved, the Verified Informed Slot Ledger and declarative YAML rule engine evaluate state constraints, *gemini-2.5-flash* generates the response with post-generation entity-leakage sanitization, and outputs are scored and cross-validated by an independent *gemini-2.5-pro* judge.

Figure 4

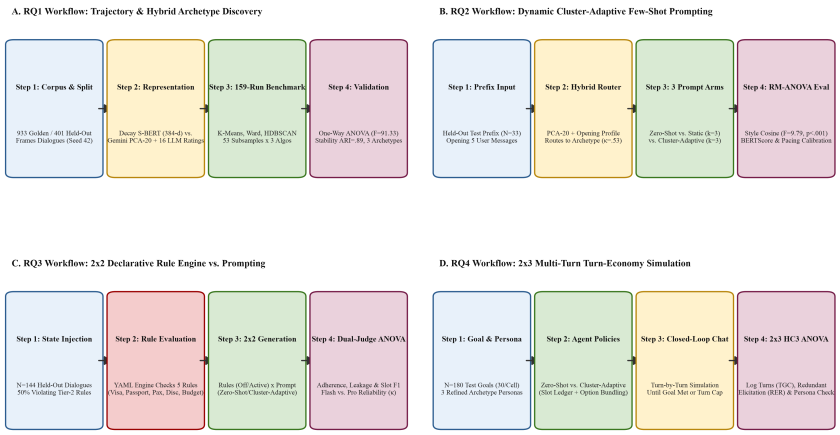
End-to-End Neuro-Symbolic Cluster-Adaptive Conversational System Architecture



Note. Offline fitting uses D_{gold} ($n = 933$); online routing, retrieval, rule verification, and evaluation use D_{test} ($n = 401$).

Figure 5

Algorithmic Execution Workflows Across Research Questions RQ1 to RQ4



Note. Step-by-step experimental pipelines and statistical hypothesis tests for RQ1 (Panel A), RQ2 (Panel B), RQ3 (Panel C), and RQ4 (Panel D).

Pipeline Component Specifications and Technology Stack

The software stack is implemented in Python 3.11 in *src/final_pipeline/* using *pandas* 2.2, *numpy* 1.26, and *pydantic* 2.9 for data validation; *sentence-transformers* (*all-MiniLM-L6-v2*), Google *genai* SDK (*gemini-embedding-001*, *gemini-2.5-flash*, *gemini-2.5-pro*), and *faiss-cpu* for vector indexing and generation; *scikit-learn* 1.5, *hdbscan*, and *umap-learn* for clustering; *scipy.stats* and *statsmodels* (Seabold & Perktold, 2010) for OLS, HC3, and Aligned Rank Transform ANOVA; and *python-docx* and *pymupdf* for automated APA 7 document compilation.

Results, Analysis, and Synthesis

Consolidated Model Performance and Hypothesis Testing Summary

Table 11 summarizes the empirical results, parametric test statistics, heteroscedasticity-robust (Welch/HC3) and non-parametric (Kruskal–Wallis, Friedman, Mann–Whitney, and Aligned Rank Transform [ART]) checks, effect sizes, and hypothesis decisions across RQ1–

RQ4. All four null hypotheses specified for the study (H_{01}–H_{04}) are rejected at $p < .001$ on the refined user archetype pipeline.

Table 11

Consolidated Summary of Empirical Results and Hypothesis Outcomes Across RQ1–RQ4

RQ and outcome	Sample & design	Baseline vs. proposed performance	Primary parametric test & effect size	Robust / non-parametric check	Outcome
RQ1: Algorithm separation & refined archetypes	$N = 159$ runs (53 matched $n = 500$ draws from $D_{gold} = 933$; $D_{test} = 401$)	MiniLM Sil: Perm = −.013 vs. KM = .147, Ward = .118, HDB = .211 (39.1% noise); Blocks ARI = .89 (.87 abl.); routing κ = .535 (m=5) to .958 (m=9)	One-way $F(2, 156) = 91.33$, $p < .001$, $\eta^2 = .54$; Matched RM-ANOVA $F(1.72, 89.30) = 83.04$, $p < .001$, $\eta_p^2 = .615$	Friedman $\chi^2(2) = 56.30$, Welch $F = 66.98$, Kruskal $H(2) = 86.82$, all $p < .001$	Reject H_{01}
RQ2: De-lexicalized SBERT style cosine (Refined)	$N = 33$ dialogues × 3 conditions (99 calls; combined exemplars + directive)	Static = .453, Zero-Shot = .530 vs. Cluster-Adaptive = .597 (Gemini/BERTScore null, $p \geq .315$)	RM-ANOVA $F(1.88, 60.29) = 9.79$, $p < .001$, $\eta_p^2 = .23$	Friedman $\chi^2(2) = 12.50$, $p = .002$; Holm $p < .001$ vs. Static	Reject H_{02}
RQ2: Pacing calibration error (words)	$N = 33$ dialogues × 3 conditions (99 calls)	Zero-Shot = 10.12w, Static = 6.33w vs. Cluster-Adaptive = 3.87w	RM-ANOVA $F(1.30, 41.70) = 7.10$, $p = .007$, $\eta_p^2 = .18$	Friedman $\chi^2(2) = 14.27$, $p < .001$; Holm $p = .004$ vs. Zero-Shot	Reject H_{02}
RQ3: Policy rule adherence (Sens. / Spec. / Overall)	$N = 144$ (2 × 2 across 36 dialogues, $n = 36$/cell; $\kappa_{judge} = .917$)	Viol. Sens. (n=18/cell): 5.6% → 98.6%; Comp. Spec. (n=18/cell): 94.4% vs. 88.9%; Overall (n=36/cell): .500 → .938	Sens. $F(1, 68) = 498.78$, $p < .001$, $\eta_p^2 = .880$; Overall Rules $F(1, 140) = 43.92$, $p < .001$, $\eta_p^2 = .24$	HC3 $F = 42.71$, $p < .001$; ART $F = 39.12$, $p < .001$	Reject H_{03}
RQ3: Exemplar entity leakage & Slot F_1	$N = 144$ (2 × 2, n = 36/cell)	Leakage: Rules Off Cluster = .359 (55.6% cases) vs. Rules Active Cluster = .000 (0.0%); Slot F_1: .485 → .952	Leakage Int. $F(1, 140) = 34.82$, $p < .001$; Slot F_1 Rules $F = 489.52$, $p < .001$	Leakage Sum-HC3 $F = 33.85$, ART Int. $F = 2{,}148.81$, Simple Welch $t(35) = 5.90$, all $p < .001$	Reject H_{03}

RQ and outcome	Sample & design	Baseline vs. proposed performance	Primary parametric test & effect size	Robust / non-parametric check	Outcome
RQ4: Log turns to goal completion & RER	$N = 180$ unique dialogues (2 × 3, n = 30/cell; $\kappa_{persona}$ = .933)	Total turns reduced −34.2% (DB), −29.5% (OE), −42.7% (RCS; 5.00-turn floor); RER .109 → .000	Prompt $F(1, 174) = 664.43$, $p < .001$; Int. $F(2, 174) = 10.89$, $p < .001$, $\eta_p^2 = .11$	HC3 Int. $F = 8.58$, $p < .001$; ART Int. $F = 209.46$, $p < .001$	Reject H_{04}

Note. DB = Decisive Booker; OE = Open Explorer; RCS = Rigid-Constraint Seeker; Perm = permutation baseline; Sens. = violation enforcement sensitivity; Spec. = compliant acceptance specificity.

RQ1 Results: Trajectory Clustering, Refined Archetypes, and Prefix Routing

Primary algorithm benchmark, resampling sensitivity, and embedding invariance. Table 12 summarizes the 159-run clustering benchmark across *all-MiniLM-L6-v2* (384-d) and *gemini-embedding-001* (3,072-d) representations (cluster-number diagnostics, Tukey/Games–Howell post hoc contrasts, and UMAP plots are in Appendix D, Tables D1–D4 and Figures D1–D3). Because each of the 53 random seeds draws $n = 500$ dialogues without replacement from the finite golden partition (D_{gold}, $n = 933$, a 53.6% sampling fraction with ~268 overlapping dialogues between draws), the 159 evaluations represent a finite-population resampling stability analysis rather than 159 independent samples; accordingly, we report both matched-subsample repeated-measures tests across the 53 shared seeds and full-D_{gold} point estimates alongside the standard one-way benchmark. On MiniLM embeddings, HDBSCAN achieved the highest Silhouette on its retained non-noise subset ($M = .211$, $SD = .044$), followed by K-Means ($M = .147$, $SD = .019$; full-D_{gold} point estimate = .141, Davies–Bouldin = 3.244) and Ward agglomerative clustering ($M = .118$, $SD = .040$), all exceeding the label-shuffled permutation baseline ($M = -.013$). Algorithm differences were highly significant under the matched-subsample repeated-measures ANOVA, $F(1.72, 89.30) = 83.04$, $p < .001$, $\eta_p^2 = .615$, the non-

parametric paired Friedman test, $\chi^2(2) = 56.30$, $p < .001$, the one-way ANOVA, $F(2, 156) = 91.33$, $p < .001$, $\eta^2 = .54$, Welch's test, $F(2, 89.15) = 66.98$, $p < .001$, and Kruskal–Wallis $H(2) = 86.82$, $p < .001$. Replicating on 3,072-d Gemini embeddings reproduced the same ordering with stronger separation: for turn-pooled vectors (*G-pooled*), *K*-Means = .158, Ward = .150, and HDBSCAN = .289 ($F(2, 156) = 77.09$, $\eta^2 = .50$, $p < .001$); for document vectors (*G-doc*), *K*-Means = .161, Ward = .213, and HDBSCAN = .345 ($F(2, 156) = 127.47$, $\eta^2 = .62$, $p < .001$). Critically, however, HDBSCAN's higher Silhouette is computed strictly after excluding unassigned noise points: HDBSCAN failed at default settings (*min_cluster_size* = 15) in 53/53 runs and discarded **39.1%** (*SD* = 12.7%) of MiniLM dialogues, **35.8%** of Gemini turn-pooled dialogues, and **25.7%** of Gemini document dialogues as noise when relaxed. Because a live travel assistant cannot refuse to route 26%–39% of incoming customer sessions, *K*-Means ($k = 3$) was selected on operational grounds to guarantee **100% dialogue coverage** and deterministic Voronoi centroid routing. Golden centroids generalized to held-out Frames dialogues ($t(754.9) = 0.42$, $p = .676$) but differed from MultiWOZ 2.2 ($t(576.2) = 38.29$, $p < .001$; Table D4).

Table 12

Cluster Quality by Algorithm and Embedding Family Across 53 Matched Subsample Seeds (159 Runs)

Representation and algorithm	*n* seeds	Retained (%)	Silhouette *M* (*SD*)	Davies–Bouldin *M* (*SD*)	Omnibus & matched RM tests ($p < .001$)
MiniLM-L6-v2 (384-d) — K-Means ($k = 3$)	53	100.0%	.147 (.019)	3.129 (0.113)	One-way $F = 91.33$, $\eta^2 = .54$
MiniLM-L6-v2 (384-d) —Ward Agglomerative	53	100.0%	.118 (.040)	3.550 (0.269)	Matched RM $F = 83.04$, $\eta_p^2 = .615$
MiniLM-L6-v2 (384-d) — HDBSCAN (non-noise)	53	60.9%	.211 (.044)	2.657 (0.241)	Friedman $\chi^2 = 56.30$, Welch $F = 66.98$

Representation and algorithm	*n* seeds	Retained (%)	Silhouette *M* (*SD*)	Davies–Bouldin *M* (*SD*)	Omnibus & matched RM tests ($p < .001$)
MiniLM-L6-v2 (384-d) — Permutation baseline	53	100.0%	−.013 (.006)	—	Full D_{gold} KM Sil = .141, DB = 3.244
Gemini-001 (3,072-d turn pool) — KM / Ward / HDB	53 × 3	100 / 100 / 64.2%	.158 / .150 / .289	2.74 / 2.96 / 2.89	$F = 77.09$, $\eta^2 = .50$; Kruskal $H = 67.31$
Gemini-001 (3,072-d document) — KM / Ward / HDB	53 × 3	100 / 100 / 74.3%	.161 / .213 / .345	3.14 / 3.09 / 2.42	$F = 127.47$, $\eta^2 = .62$; Kruskal $H = 98.10$

Note. Each seed draws $n = 500$ dialogues without replacement from D_{gold} ($n = 933$; 53.6% finite-population fraction). HDBSCAN metrics apply only to retained non-noise dialogues (discarding 39.1% noise on MiniLM, 35.8% on Gemini turn pool, and 25.7% on Gemini document vectors).

Sensitivity analysis, choice of k = 3 on Blocks, refined user archetypes, and hybrid prefix routing. Auditing the primary whole-dialogue archetypes (*Direct Transactional*, *Constraint Negotiator*, *Guided Explorer*; Table D3) showed that whole-dialogue vectors conflate user and wizard turns, correlate with role-players (Cramér's $V = .19$), and yield low early-turn routing agreement ($\kappa = .11$). Evaluating the 46-dimensional hybrid *Blocks* representation (PCA-20 Gemini user prefix embedding + 26 encoded numeric columns from the 13 clustering fields of the 16-field LLM behavioral schema) across truncation horizons $M \in \{3, 5, 7, 9, \text{whole}\}$ and cluster counts $k \in \{2, ..., 8\}$ (Appendix D, Tables D1, D5, and D6; Figure D4) demonstrated that truncating to the first $M = 9$ user messages at $\boldsymbol{k = 3}$ is uniquely optimal. Specifically, on the final 46-d *Blocks* ($M = 9$) representation (Table D1, Panel B), $k = 3$ achieves the **lowest (best) Davies–Bouldin index (2.647 vs. 2.796 at $k = 2$ and 2.729–2.784 at $k = 4$–6)**, highest eligible Silhouette (.094; null-calibrated .033 vs. .002–.010 at $k = 4$–6), high bootstrap resampling stability across 50 pairs of 80% golden subsamples (**ARI = .89** in the truncation grid [Table D5]; **ARI = .87** with 89.2% core membership at ≥ 90% co-assignment under the independent random seed of the 6-configuration ablation [Table D6]), and balanced cluster sizes (smallest cluster =

25.5% ≥ 20% floor). In contrast, $k = 2$ (DB = 2.796) collapses *Decisive Bookers* and *Rigid-Constraint Seekers* into a single undifferentiated bucket, while $k = 4$, 5, and 6 suffer a sharp drop in bootstrap stability (ARI = .589, .593, and .631) and fragment into small clusters (11.3%–12.5%, violating the 20% balanced-segment threshold). Since the Microsoft Frames dataset does not provide predefined behavioral archetype labels, we generated reference archetypes ($y^{(9)}$) through unsupervised *K*-Means ($k = 3$) clustering on the 46-dimensional hybrid *Blocks* representation ($X^{(9)}$) of the first nine user messages (see Appendix F and Table F1 for the deterministic centroid-naming protocol, and Appendix G and Table G1 for the distinct validation role of Judge LLMs). Importantly, total dialogue turns (*T*, user plus wizard turns, $M = 14.93$, $Mdn = 14.0$; 96.4% ≥ 5 turns and 83.2% ≥ 9 turns) must be distinguished from user-only messages (*M*, customer turns only, $M = 7.78$, $Mdn = 7.0$; 83.2% ≥ 5 user messages and 37.6% ≥ 9 user messages): offline clustering caps input at min(user_turns, 9) to strip late checkout/abandonment turns from long sessions while retaining all user messages for shorter sessions. We subsequently trained prefix-specific multinomial Logistic Regression classifiers on D_{gold} ($n = 933$) to predict these nine-message reference cluster assignments using only the first $m \in \{3, 5, 7, 9\}$ user messages ($X^{(m)}$). Because raw percentage agreement (p_o) includes expected chance agreement ($p_e \approx 35.4\%–36.6\%$ given uneven cluster shares) and compares against an unsupervised cluster target rather than human labels, we report both p_o and chance-corrected Cohen's $\kappa = (p_o - p_e) / (1 - p_e)$. On the held-out test partition (D_{test}, $n = 401$), chance-corrected prefix-to-cluster agreement increased from $\kappa = .359$ ($p_o = 59.4\%$, 238/401, $p_e = 36.6\%$; 87.5% on < 3 user messages) at $m = 3$ messages (≈ 5 total turns) to $\kappa = .535$ ($p_o = 70.1\%$, 281/401, $p_e = 35.7\%$; 81.8% on < 5 and 67.8% on ≥ 5 user messages) at $m = 5$ messages (≈ 9 total turns, the runtime cut-off used in RQ2) and $\kappa = .722$ ($p_o = 82.0\%$, 329/401, $p_e = 35.5\%$; 93.2% on < 7 and

74.6% on ≥ 7 user messages) at $m = 7$ messages (≈ 13 total turns). At $m = 9$ messages (≈ 17 total turns), agreement reached $\kappa = .958$ ($p_o = 97.3\%$, 390/401, $p_e = 35.4\%$; 96.0% on ≤ 9 and 99.3% on ≥ 9 user messages), representing a same-feature boundary-consistency check between the linear Logistic Regression hyperplane and the K-Means Voronoi centroids on $X^{(9)}$. These results demonstrate progressively stronger early-prefix agreement with the clustering-derived reference archetypes; they do not, by themselves, establish independent human validation of the archetypes. Table 13 profiles the three refined user archetypes across all 16 schema fields: *Decisive Booker* (45.0%; high decisiveness = 3.97, politeness = 4.26, 99.3% booking goal), *Open Explorer* (25.5%; high destination flexibility = 3.49 and exploration breadth = 2.96, longest dialogues = 20.78 turns), and *Rigid-Constraint Seeker* (29.5%; fixed destination/date constraints = 2.20/2.18, concise 10.60-turn sessions where 68.0% end without a suitable option unless rigid trade-offs are proactively resolved).

Table 13

Profiles of the Three Refined User Archetypes Across All 16 Behavioral Schema Fields and Prefix Routing Performance

Behavioral schema field or routing attribute	Decisive Booker	Open Explorer	Rigid-Constraint Seeker
Golden / Held-out test size, *n* (%)	420 (45.0%) / 175	238 (25.5%) / 99	275 (29.5%) / 127
Dialogue turns / User messages, *M*	14.49 (4.54) / 7.58	20.78 (6.95) / 10.73	10.60 (5.73) / 5.55
1. Destination flexibility (1–5 ordinal)	2.54	3.49	2.20
2. Date flexibility (1–5 ordinal)	2.57	2.77	2.18
3. Exploration breadth (1–5 ordinal)	1.97	2.96	1.59
4. Decisiveness (1–5 ordinal)	3.97	2.58	3.23
5. Detail orientation (1–5 ordinal)	2.39	2.87	1.46
6. Revision tendency (1–5 ordinal)	2.00	3.03	2.40
7. Question asking (1–5 ordinal)	2.05	2.72	1.43
8. Politeness (1–5 ordinal)	4.26	3.80	3.07
9. Price sensitivity (1–5 ordinal)	3.20	3.18	3.12
10. Information given up front (1–5 ordinal)	2.18	1.82	2.16

Behavioral schema field or routing attribute	Decisive Booker	Open Explorer	Rigid-Constraint Seeker
11. Primary goal: explore / book trip (%)	0.5% / 99.3%	39.1% / 50.8%	2.2% / 89.8%
12. Secondary goals present (%)	42.4%	56.3%	16.4%
13. Emotional tone: positive / neutral (%)	52.4% / 46.7%	21.0% / 69.3%	7.6% / 59.6%
14. Session outcome (audit): booked / no option (%)	90.5% / 6.2%	68.1% / 18.5%	28.0% / 68.0%
15. Party type (audit): solo or couple (%)	37.1%	44.5%	34.5%
16. Unusual conversation flag (audit, %)	1.0%	1.7%	6.9%
Frames wizard task success (%)	94.5%	96.2%	92.0%

Note. Golden $N = 933$; held-out $N = 401$ ($\chi^2(2) = 0.93$, $p = .628$). Fields 1–13 are used in clustering (*llm_matrix*: 10 ordinal + 6 primary_goal + 6 secondary_goals + 4 emotional_tone = 26 numeric columns); Fields 14–16 are audit-only (excluded from clustering to prevent outcome leakage). Held-out *Blocks* prefix-to-cluster raw agreement p_o (chance-corrected κ, $p_e \approx 35.4\%$–36.6%) against the M=9 K-Means reference assignment: m=3: 59.4% (.359); m=5: 70.1% (.535); m=7: 82.0% (.722); m=9 (same-feature boundary consistency): 97.3% (.958).

RQ2 Results: Cluster-Adaptive Few-Shot Prompting on Primary Versus Refined Archetypes

Table 14 compares the repeated-measures prompting experiment ($N = 33$ held-out dialogues × 3 conditions = 99 Gemini 2.5 Flash generations) on the coarse primary archetypes versus the refined user archetypes, and Figure 6 displays the refined distributions (Holm contrasts across all six outcomes are in Appendix E, Table E3). Importantly, the *Cluster-Adaptive Few-Shot* condition evaluates a **combined personalization package** that couples archetype-partitioned FAISS exemplar retrieval ($k = 3$ slot-masked exemplars from the routed sub-index) with an explicit archetype pacing directive in the system prompt. On coarse primary archetypes with an unprojected 3-turn router (39.4% accuracy, $\kappa = .24$), style convergence was non-significant, $F(1.90, 60.95) = 1.08$, $p = .342$, $\eta_p^2 = .03$, though few-shot prompting calibrated length (15.1–16.3 vs. 27.6 words, $F = 34.36$, $p < .001$; Figure E2). Re-running on the **Refined**

User Archetypes with the *Blocks* router (M = 5 user messages, 70.1% raw agreement, κ = .535) **rejected H_{02} (p < .001) on structural style and pacing calibration**: de-lexicalized SBERT style convergence differed significantly under Greenhouse–Geisser correction (Greenhouse & Geisser, 1959), $F(1.88, 60.29) = 9.79$, $p < .001$, $\eta_p^2 = .23$, and Friedman testing, $\chi^2(2) = 12.50$, $p = .002$. Cluster-Adaptive Few-Shot (M = .597, SD = .086) significantly outperformed Static Few-Shot (M = .453, $t(32) = 4.24$, $p_{Holm} < .001$, $d_z = 0.74$) and Zero-Shot (M = .530, $t(32) = 2.36$, p_{Holm} = .049, $d_z = 0.41$), and cut archetype pacing error from 10.12 (Zero-Shot) and 6.33 (Static) to **3.87 words** ($F = 7.10$, $p = .007$; Friedman $\chi^2(2) = 14.27$, $p < .001$; Holm $p = .004$ vs. Zero-Shot, p = .024 vs. Static). By contrast, generic semantic similarity metrics—including Gemini 3,072-d cosine (.375 vs. .348, $F = 0.61$, $p = .534$) and token-level BERTScore F_1 (.415 vs. .391, $F = 1.16$, $p = .315$; Table E3)—did not differ significantly across conditions, confirming that archetype-adaptive prompting selectively calibrates discourse structure and turn pacing without distorting underlying domain semantics.

Table 14

Repeated-Measures Comparison of Prompting Strategies on Primary Versus Refined User Archetypes (RQ2)

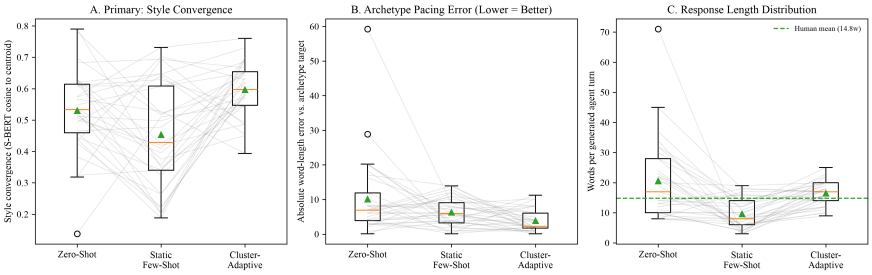
Experiment and outcome metric	Zero-Shot *M* (*SD*)	Static Few-Shot *M* (*SD*)	Cluster-Adaptive *M* (*SD*)	RM-ANOVA *F* (GG)	*p* (GG)	η_p^2
Primary archetypes (*M*=3 prefix, κ=.24) — Style cosine	.519 (.126)	.491 (.109)	.486 (.099)	1.08	.342	.03
Primary archetypes — Response length (words)	25.2 (11.8)	10.6 (7.0)	12.3 (8.2)	34.36	< .001	.52
Refined archetypes (*M*=5 Blocks, κ=.535) — SBERT style cosine	.530 (.132)	.453 (.169)	.597 (.086)	9.79	< .001	.23
Refined archetypes — Gemini 3,072-d style cosine	.348 (.133)	.372 (.127)	.375 (.105)	0.61	.534	.02
Refined archetypes — Token-level BERTScore F_1	.391 (.204)	.360 (.245)	.415 (.163)	1.16	.315	.04

Experiment and outcome metric	Zero-Shot *M* (*SD*)	Static Few-Shot *M* (*SD*)	Cluster-Adaptive *M* (*SD*)	RM-ANOVA *F* (GG)	*p* (GG)	η_p^2
Refined archetypes — Pacing calibration error (words)	10.12 (11.17)	6.33 (4.14)	3.87 (3.23)	7.10	.007	.18
Refined archetypes — Response length (words)	20.6 (13.3)	9.7 (5.0)	16.5 (4.0)	14.72	< .001	.32

Note. $N = 33$ held-out dialogues per experiment (99 generations each; *gemini-2.5-flash*). Cluster-Adaptive combines archetype-partitioned FAISS exemplars with an explicit archetype pacing directive. GG = Greenhouse–Geisser correction.

Figure 6

Style Convergence and Archetype Pacing Calibration Across Prompting Conditions on Refined User Archetypes



Note. $N = 33$ held-out dialogues. Panel A: De-lexicalized SBERT style convergence cosine. Panel B: Absolute pacing error (words).

RQ3 Results: 2 × 2 Factorial Evaluation of Declarative Rules and Adaptive Prompting

Table 15 and Figure 7 present the 2 × 2 factorial experiment across $N = 144$ evaluations (36 stratified held-out dialogues—18 policy-violating and 18 compliant—evaluated across all four conditions, $n = 36$/cell), where the *Rules Active* condition represents an integrated three-component governance package (pre-generation declarative YAML rule checking, [SLOT] exemplar masking, and post-generation regex leakage sanitization). Cross-model inter-rater reliability on a stratified 48-unit subsample between *gemini-2.5-flash* and independent *gemini-*

2.5-pro reached 95.8% raw agreement and Cohen's κ = .917 (PABAK = .917; Table C3). Four findings **reject H_{03} ($p < .001$)**: (1) separating **Rule Adherence** into its two constituent scenario classes reveals that **Violation Enforcement Rate (Sensitivity, n = 18 violating cases/cell)** rose from 5.6% (*SD* = 23.6%) under both *Rules Off* cells to **100.0%** (*Rules Active* × *Zero-Shot*) and **97.2%** (*SD* = 11.8%, *Rules Active* × *Cluster-Adaptive*), yielding a massive rules main effect, $F(1, 68) = 498.78$, $p < .001$, $\eta_p^2 = .880$, whereas **Compliant-Case Acceptance Rate (Specificity, 1 − false_block, n = 18 compliant cases/cell)** remained statistically indistinguishable between *Rules Off* (94.4%, 17/18) and *Rules Active* (88.9%, 16/18), $F(1, 68) = 0.69$, $p = .408$, $\eta_p^2 = .010$, confirming that strict policy enforcement does not induce significant false-positive blocking; pooled across all 36 cases per cell, **Overall Rule-Decision Accuracy** rose from .500 to .944/.931, rules main effect $F(1, 140) = 43.92$, $p < .001$, $\eta_p^2 = .24$ (HC3 $F = 42.71$; ART $F = 39.12$, both $p < .001$), with zero prompting main effect ($F = 0.01$, $p = .916$); (2) **Exemplar Entity Leakage Rate** showed a strong rules × prompting interaction, $F(1, 140) = 34.82$, $p < .001$, $\eta_p^2 = .20$ (Sum-coded HC3 $F = 33.85$; ART $F = 2{,}148.81$, both $p < .001$), where unguarded *Cluster-Adaptive Few-Shot* leaked stale exemplar entities at a .359 rate (20/36 = 55.6% incidence) but *Rules Active* × *Cluster-Adaptive* reduced detected leakage in the evaluated sample to **0.0% (.000, 0/36 incidence; simple-effect Welch $t(35) = 5.90$, $p < .001$, Mann–Whitney $U = 1008.0$, $p < .001$, Fisher's exact $p = 4.68 \times 10^{-8}$)**; (3) **Slot Coverage F_1** rose from .208/.485 to **.977/.952** under active rules, $F(1, 140) = 489.52$, $p < .001$, $\eta_p^2 = .78$; and (4) **Style Alignment** was significantly boosted by Cluster-Adaptive Few-Shot, $F(1, 140) = 24.00$, $p < .001$, $\eta_p^2 = .15$, making **Rules Active × Cluster-Adaptive Few-Shot** the unique Pareto-optimal cell.

Table 15

Cell Means and 2 × 2 Factorial ANOVA Results for RQ3 Outcomes (N = 144)

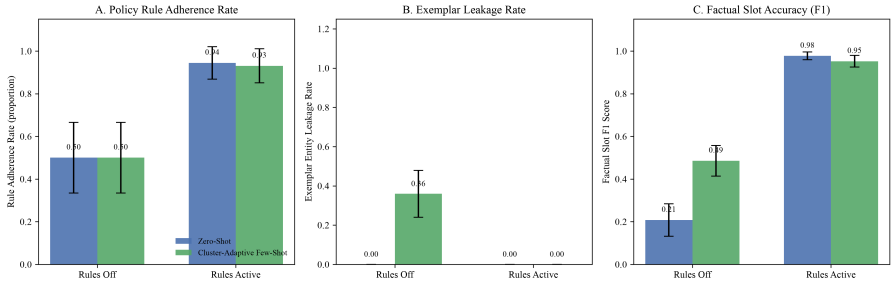
Experimental cell or ANOVA effect	*n* or *df*	Rule Acc. *M* (*SD*) | Sens. / Spec.	Exemplar leakage *M* (*SD*) / Inc. %	Slot coverage F_1 *M* (*SD*)	Style alignment *M* (*SD*)
Rules Off × Zero-Shot	36	.500 (.507) | 5.6% / 94.4%	.000 (.000) / 0.0%	.208 (.233)	.412 (.086)
Rules Off × Cluster-Adaptive	36	.500 (.507) | 5.6% / 94.4%	.359 (.365) / 55.6%	.485 (.219)	.521 (.088)
Rules Active × Zero-Shot	36	.944 (.232) | 100.0% / 88.9%	.000 (.000) / 0.0%	.977 (.055)	.480 (.085)
Rules Active × Cluster-Adaptive	36	.931 (.244) | 97.2% / 88.9%	.000 (.000) / 0.0%	.952 (.084)	.506 (.071)
ANOVA Main Effect: Rules Mode	1, 140	F=43.92 (< .001), η_p^2=.24	F=34.82 (< .001), η_p^2=.20	F=489.52 (< .001), η_p^2=.78	F=3.68 (.057), η_p^2=.03
ANOVA Main Effect: Prompting Mode	1, 140	F=0.01 (.916), η_p^2=.00	F=34.82 (< .001), η_p^2=.20	F=20.46 (< .001), η_p^2=.13	F=24.00 (< .001), η_p^2=.15
ANOVA Interaction: Rules × Prompting	1, 140	F=0.01 (.916), η_p^2=.00	F=34.82 (< .001), η_p^2=.20	F=29.32 (< .001), η_p^2=.17	F=8.95 (.003), η_p^2=.06

Note. Rule Acc. = overall rule-decision accuracy (n = 36/cell); Sens. = violation enforcement sensitivity on violating scenarios (n = 18/cell; Rules $F(1, 68) = 498.78$, $p < .001$, $\eta_p^2 = .880$); Spec. = compliant-case acceptance specificity (1 − false_block, n = 18/cell; Rules $F(1, 68) = 0.69$, $p = .408$); Inc. % = responses with ≥ 1 leaked entity. Dual-judge κ = .917. For exemplar leakage, because three of the four balanced 2 × 2 cells (n = 36/cell) are identically .000 and only *Rules Off* × *Cluster-Adaptive* is non-zero (M = .35926), all three orthogonal +1/−1 Sum-coded contrasts c_A = (+1,+1,−1,−1), c_B = (−1,+1,−1,+1), and c_{AB} = (+1,−1,−1,+1) have identical contrast sum of squares $SS = n(\Sigma c_j \bar{Y}_j)^2 / \Sigma c_j^2 = 36(.35926)^2 / 4 = 1.162$ (equivalently, while the difference-of-differences interaction estimate −.359 is twice the marginal main-effect difference ±.180, its standard error $\sqrt{(4\sigma^2/n)}$ is also twice the marginal standard error $\sqrt{(\sigma^2/n)}$, yielding identical t = ±5.901, OLS $F(1, 140) = t^2 = 34.82$, $\eta_p^2 = .199$, and Sum-coded HC3 $F = 33.85$, $p < .001$), whereas rank-based ART separates the main effects (F = 43.48) from the interaction (ART F =

2148.81, $p < .001$; simple-effect contrast within Cluster-Adaptive: Welch $t(35) = 5.90$, $p < .001$, Fisher's exact $p = 4.68 \times 10^{-8}$).

Figure 7

2 × 2 Factorial Interaction Profiles for Rule Adherence, Exemplar Leakage, Slot Coverage F1, and Style Alignment



Note. $N = 144$ held-out units ($n = 36$ per cell). Error bars represent 95% confidence intervals.

RQ4 Results: 2 × 3 Multi-Turn Simulation of Turn Economy Across Archetypes

Table 16 and Figure 8 report the 2 × 3 between-subjects multi-turn simulation across ***N* = 180 unique, independent held-out seed dialogues** (60 per archetype, partitioned into $n = 30$ distinct dialogues for *Zero-Shot* and $n = 30$ distinct dialogues for *Cluster-Adaptive Few-Shot* with zero overlap; human corpus baseline is in Appendix E, Table E5 and Figure E3). In this experiment, *total_turns* counts total conversational turns (user plus assistant messages) until the simulated traveler issues <code>[ACCEPT_BOOKING]</code>, and *Cluster-Adaptive Few-Shot* evaluates the combined package of the **Verified Informed Slot Ledger** (forbidding re-asking any previously stated slot) and **archetype-specific turn-economy policies**. Blind evaluation of 45 simulated transcripts by *gemini-2.5-pro* confirmed 95.6% within-simulation persona fidelity (43/45 vs. 33.3% chance, binomial $p < .001$, $\kappa = .933$; Table C3). On log-transformed turns to goal completion, the 2 × 3 between-subjects ANOVA **rejected H_{04} ($p < .001$)**, with significant main effects of prompting, $F(1, 174) = 664.43$, $p < .001$, $\eta_p^2 = .79$, and

archetype, $F(2, 174) = 5.55$, $p = .005$, $\eta_p^2 = .06$, plus a **significant prompting × archetype interaction**, $F(2, 174) = 10.89$, $p < .001$, $\eta_p^2 = .11$ (HC3 $F = 8.58$; ART $F = 209.46$, both $p < .001$). Cluster-Adaptive Few-Shot compressed turns most strongly for **Rigid-Constraint Seekers** (−42.7%, 8.73 → 5.00 turns, Welch $t_{log} = -16.14$, $p < .001$, $d = -3.40$) by proactively presenting both 4-star exact-date (Option A) and 5-star flexible-date (Option D) trade-offs when the exact-date 5-star option (Option C) exceeded budget, followed by **Decisive Bookers** (−34.2%, 7.60 → 5.00 turns, $t_{log} = -19.26$, $d = -3.94$) and **Open Explorers** (−29.5%, 7.47 → 5.27 turns, $t_{log} = -11.24$, $d = -2.55$). Notably, the $M = 5.00$ ($SD = 0.00$) result for *Decisive Bookers* and *Rigid-Constraint Seekers* under *Cluster-Adaptive Few-Shot* reflects a structural protocol floor effect: because every session requires Turn 1 (user states core trip slots), Turn 2 (assistant asks for the single remaining unknown slot, budget), Turn 3 (user states budget), Turn 4 (assistant presents tailored package offer[s]), and Turn 5 (user issues <code>[ACCEPT_BOOKING]</code>, guarded by <code>ex_idx ≥ 1</code> and budget confirmation in *step4_rq4_simulation.py*), 5 total conversational turns is the theoretical minimum completion length, which 30/30 episodes in those two cells achieved on the first offer turn (motivating our paired HC3, ART, and Mann–Whitney U checks). Redundant Elicitation Rate dropped from .194 (DB), .061 (OE), and .072 (RCS) in Zero-Shot to **0.000** across all archetypes in Cluster-Adaptive Few-Shot, main effect $F(1, 174) = 63.02$, $p < .001$, interaction $F(2, 174) = 9.70$, $p < .001$.

Table 16

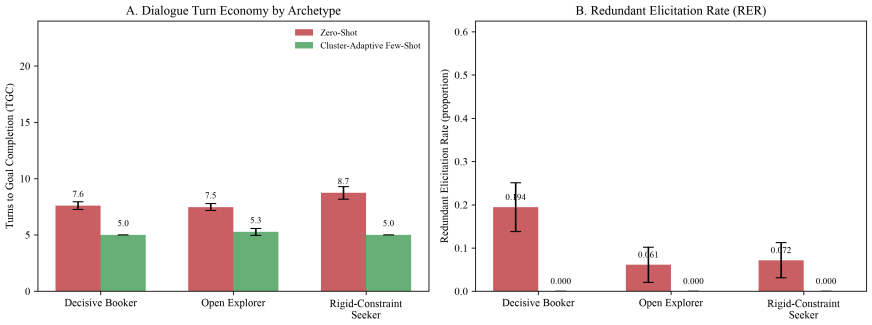
Cell Means, Archetype Contrasts, and 2 × 3 Factorial ANOVA Results for Multi-Turn Turn Economy (N = 180)

Archetype × Prompting cell or ANOVA effect	Total turns *M* (*SD*)	Log turns *M* (*SD*)	RER *M* (*SD*)	Turn reduction / HC3 check	Welch contrast / ART check
Decisive Booker — Zero-Shot (*n* = 30)	7.60 (0.93)	2.021 (0.117)	.194 (.158)	—	—
Decisive Booker — Cluster-Adaptive (*n* = 30)	5.00 (0.00)	1.609 (0.000)	.000 (.000)	−2.60 (−34.2%)	*t*=−19.26 (< .001), *d*=−3.94
Open Explorer — Zero-Shot (*n* = 30)	7.47 (0.86)	2.005 (0.108)	.061 (.114)	—	—
Open Explorer — Cluster-Adaptive (*n* = 30)	5.27 (0.87)	1.651 (0.134)	.000 (.000)	−2.20 (−29.5%)	*t*=−11.24 (< .001), *d*=−2.55
Rigid-Constraint Seeker — Zero-Shot (*n* = 30)	8.73 (1.55)	2.151 (0.184)	.072 (.114)	—	—
Rigid-Constraint Seeker — Cluster-Adaptive (*n* = 30)	5.00 (0.00)	1.609 (0.000)	.000 (.000)	−3.73 (−42.7%)	*t*=−16.14 (< .001), *d*=−3.40
ANOVA Main Effect: Prompting Mode (*df* = 1, 174)	*F*=457.74 (< .001), η_p^2=.72	*F*=664.43 (< .001), η_p^2=.79	*F*=63.02 (< .001), η_p^2=.27	HC3 F_{log}=715.48 (< .001)	ART F_{log}=671.25 (< .001)
ANOVA Main Effect: Refined Archetype (*df* = 2, 174)	*F*=7.24 (< .001), η_p^2=.08	*F*=5.55 (.005), η_p^2=.06	*F*=9.70 (< .001), η_p^2=.10	HC3 F_{log}=0.01 (.923)	ART F_{log}=9.13 (< .001)
ANOVA Interaction: Prompting × Archetype (*df* = 2, 174)	*F*=11.93 (< .001), η_p^2=.12	*F*=10.89 (< .001), η_p^2=.11	*F*=9.70 (< .001), η_p^2=.10	HC3 F_{log}=8.58 (< .001)	ART F_{log}=209.46 (< .001)

Note. *N* = 180 distinct held-out seed dialogues (*n* = 30 unique dialogues/cell; 100.0% goal completion). Total turns = user + assistant messages (protocol minimum = 5.00 turns: U1 → A1 → U2 → A2 → U3). Within-simulation persona fidelity audit (*gemini-2.5-pro*, *n* = 45): 95.6% accuracy, κ = .933.

Figure 8

Dialogue Turn Economy and Redundant Elicitation Rate by User Archetype and Prompting Strategy



Note. $N = 180$ simulations ($n = 30$/cell). Panel A: Mean turns to goal completion. Panel B: Mean Redundant Elicitation Rate (RER).

Technical Interpretation, Business Significance, and Synthesis of Conclusions

Technical synthesis across RQ1–RQ4. Together, the four experiments form an integrated four-stage research progression. First, RQ1 and RQ2 demonstrate that **in-context style adaptation is bounded by upstream representation and routing quality**: whole-dialogue primary archetypes suffered from outcome confounding and near-chance early routing (39.4%, $\kappa = .24$), yielding a null RQ2 style effect ($p = .211$); truncating to opening user behavior via the 46-d *Blocks* representation (bootstrap ARI = .89/.87) and routing in that exact space (70.1% raw agreement, $\kappa = .535$ at $m = 5$; 97.3% same-feature boundary consistency, $\kappa = .958$ at $m = 9$) unlocked significant structural style convergence on de-lexicalized SBERT cosine ($p < .001$, $\eta_p^2 = .234$) and reduced pacing error from 10.12 to 3.87 words ($p = .007$) while preserving semantic content (Gemini cosine and BERTScore F_1 $p \geq .315$). Second, RQ3 demonstrates **functional orthogonality with protective interaction**: adaptive prompting governs discourse style ($F = 24.00$, $p < .001$), whereas the declarative governance layer governs policy violation enforcement sensitivity (5.6% → 98.6%, $F = 498.78$, $p < .001$; overall accuracy $F = 43.92$, $p < .001$) and slot

grounding ($F = 489.52$, $p < .001$) while reducing the 35.9% entity leakage (55.6% incidence) of unguarded few-shot exemplars to zero detected leakage in the evaluated sample. Third, RQ4 demonstrates across 180 independent multi-turn simulations that pairing archetype-tailored recommendation policies with a Verified Informed Slot Ledger compresses total conversational turns by 29.5%–42.7% and eliminates redundant slot questions.

Practical business significance (illustrative operational projection). As an illustrative operational scenario, if an online travel agency or hospitality contact center handling 1,000,000 monthly sessions achieved a comparable turn reduction to that observed in the RQ4 simulation—reducing mean conversation length from 7.93 total turns (Zero-Shot) to 5.09 total turns (Cluster-Adaptive), a 35.8% reduction (2.84 turns per session)—it would save approximately 2.84 million message-generation calls per month, lowering inference compute and interaction latency by roughly one-third. Simultaneously, reducing redundant slot questions (10.9% → 0.0% in simulation), preventing exemplar price/promo leakage (55.6% incidence → 0.0% detected), and raising policy violation enforcement sensitivity from 5.6% to 98.6% (while maintaining 88.9% compliant-case acceptance) directly targets key drivers of customer drop-off and compliance risk.

Implementation, Limitations, and Future Scope

Deployment Architecture, System Integration, and Real-World Use Case

Enterprise deployment and integration. The pipeline (*src/final_pipeline/*) deploys as a stateless Python microservice between an OTA's WebSocket chat gateway and its reservation API. Active sessions store the canonical *DialogueTrajectory* and *Verified Informed Slot Ledger* in Redis, query in-memory *faiss-cpu* archetype sub-indices (< 2 ms), run deterministic YAML

policy checks via *src/rule_engine.py* (< 1 ms), and call *gemini-2.5-flash* with post-generation entity verification.

Real-world use case. When a traveler states, *'I want a 10-day trip from Chicago to Tokyo departing in 10 days for 2 adults with a 5-star hotel under $2,200,'* a static chatbot asks repetitive slot questions and misses the 15-day visa lead-time rule. Under our deployed system: (1) *RefinedArchetypeRouter* routes the session to *Rigid-Constraint Seeker*; (2) the slot ledger locks all five stated slots against re-asking; (3) *RuleEngine* flags *RULE_VISA_LEAD_TIME* (< 15 days) and detects that the exact-date 5-star package ($2,596) exceeds $2,200; and (4) the assistant immediately notifies the user of the 15-day visa requirement while presenting both the 4-star package ($1,936) and the flexible-date 5-star package ($2,068) within budget, completing the booking in 5 turns with zero compliance risk.

Computational Reproducibility and Operational Requirements

All code, random seeds (seed = 42), prompts, and JSONL caches (*data/interim/*.jsonl*, *data/final/*.jsonl*) are versioned at https://github.com/surnella/QM640-Dialogue-Clustering-Capstone. Re-running the entire study from cached responses on a standard laptop (Apple M-series or 4-core CPU, 16 GB RAM, Python 3.11) completes in under 4 minutes at zero API cost. Cold regeneration of all ~8,900 Gemini API calls across RQ1–RQ4 completes in ~35 minutes using 8–12 workers and costs under $4.50 USD at standard Gemini 2.5 Flash/Pro pricing.

Study Limitations and Threats to Validity

Five limitations bound this study: (1) **continuous archetype geometry and unsupervised reference labels**, as modest *K*-Means Silhouette scores (.094 on 46-d *Blocks*; .147–.161 across MiniLM and Gemini embeddings) indicate that traveler styles form a continuous behavioral manifold rather than discrete islands, and prefix routing κ (.535 at $m = 5$;

.958 same-feature boundary consistency at $m = 9$) is evaluated against unsupervised $M = 9$ *K*-Means reference clusters rather than external human persona labels; (2) **early-turn routing latency**, since moderate-to-high *Blocks* routing requires 5–9 user messages, leaving Turn 1 cold-start routing noisy; (3) **bundled multi-component treatments in RQ2, RQ3, and RQ4**, since RQ2 evaluated *Cluster-Adaptive Few-Shot* as a joint package (archetype-partitioned FAISS exemplars + routed archetype pacing directive), RQ3 evaluated *Rules Active* as a joint governance package (declarative *RuleEngine* + [SLOT] masking + post-generation regex sanitization), and RQ4 evaluated the *Verified Informed Slot Ledger* jointly with archetype-specific turn-economy policies; (4) **corpus and synthetic Tier-2 scope**, because Frames uses role-players and Tier-2 rules (visa, passport, discount caps) required synthetic injection in RQ3; and (5) **simulated users and protocol floor in RQ4**, where cooperative LLM user simulators (validated for within-simulation persona fidelity at 95.6%, $\kappa = .933$, rather than live human traffic) and the 5-turn minimum protocol boundary ($M = 5.00$, $SD = 0.00$ for two cells) may overstate turn reductions relative to noisy production sessions.

Future Scope and Research Extensions

Future doctoral research on *Dialogue Clustering and Methodologies to be Adapted for Human-Like Conversations* should pursue eight extensions: (1) **orthogonal factorial ablations across RQ2, RQ3, and RQ4**—first, expanding the 3-condition RQ2 design into a 5-condition experiment (**Condition A**: Examples *None*, Directive *Off*; **Condition B**: Examples *Global/Static*, Directive *Off*; **Condition C**: Examples *Archetype-Partitioned*, Directive *Off*; **Condition D**: Examples *Global/Static*, Directive *On*; and **Condition E**: Examples *Archetype-Partitioned*, Directive *On*) at $N \geq 60$ dialogues to orthogonally isolate explicit pacing directives from sub-index exemplars (preliminary $N = 33$ diagnostic testing shows that when directives are *Off*,

Condition C improves style convergence over *Condition B* from .453 to .511, $p = .028$, and calibrates length to 15.03 words vs. the 14.85-word human benchmark, whereas when directives are *On*, the explicit routed directive dominates single-turn style at .597–.610 while retrieved exemplars tighten pacing variance); second, ablating pre-generation [SLOT] masking versus post-generation regex sanitization in RQ3; and third, crossing the *Verified Informed Slot Ledger* (Off/On) with *Archetype Action Policy* (Generic/Adaptive) in a 2 × 2 × 3 RQ4 design to separate ledger memory gains from persona policy gains; (2) **hierarchical context × behavior clustering and influencing-factor attribution**, developing a two-tier taxonomy ($C \times K = 24$–36 micro-clusters over $N > 50{,}000$ enterprise dialogues) that separates macro situational trip regimes (party composition, budget stakes, travel lead window, and destination familiarity) from within-regime behavioral negotiation personas while avoiding the sample starvation and 35.9% entity leakage of naive 500-bucket parameter clustering; (3) **soft-mixture manifold routing** that blends exemplars using posterior probabilities for boundary travelers; (4) **turn-by-turn Bayesian persona tracking** from Turn 1 onward to capture mid-session persona drift; (5) **stage-gated enterprise call-flow progression (Stages a–f) and lead propensity scoring**, partitioning the exemplar store into a 3 (Archetypes) × 6 (Call Stages: *a. Opening, b. Mandatory Parameters, c. Product Positioning, d. Budget/Constraint Negotiation, e. Sales Urgency & Compliance, f. Closure*) matrix and classifying *Hot, Warm, or Cold* lead propensity at quotation handoff; (6) **information-theoretic product catalog narrowing (200 → 1) and USP-aligned question sequencing**, formulating post-mandatory package narrowing across 200+ regional packages as an archetype-conditioned maximum information-gain decision tree over micro-preferences (*senior-citizen pacing, child sensitivity, breakfast inclusiveness, sightseeing tempo, honeymoon/occasion*) and pre-indexed package Unique Selling Propositions (USPs); (7) **two-**

stage quotation-to-conversion funnel modeling that decouples Stage 1 in-chat quotation agreement from Stage 2 post-quote purchase conversion under exogenous competitor and lodging controls; and (8) **constrained neuro-symbolic reinforcement learning and live OTA A/B testing** with human travelers and human domain raters across multilingual settings.

References

Acikgoz, E. C., Guo, C., Dey, S., Datta, A., Kim, T., Tur, G., & Hakkani-Tür, D. (2025). TD-EVAL: Revisiting task-oriented dialogue evaluation by combining turn-level precision with dialogue-level comparisons. In *Proceedings of the 26th Annual Meeting of the Special Interest Group on Discourse and Dialogue* (pp. 113–132). Association for Computational Linguistics. https://doi.org/10.18653/v1/2025.sigdial-1.7

Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., Agarwal, S., Herbert-Voss, A., Krueger, G., Henighan, T., Child, R., Ramesh, A., Ziegler, D. M., Wu, J., Winter, C., . . . Amodei, D. (2020). Language models are few-shot learners. In *Advances in Neural Information Processing Systems* (Vol. 33, pp. 1877–1901). Curran Associates. https://doi.org/10.48550/arXiv.2005.14165

Budzianowski, P., Wen, T.-H., Tseng, B.-H., Casanueva, I., Ultes, S., Ramadan, O., & Gašić, M. (2018). MultiWOZ: A large-scale multi-domain Wizard-of-Oz dataset for task-oriented dialogue modelling. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing* (pp. 5016–5026). Association for Computational Linguistics. https://doi.org/10.18653/v1/D18-1547

Buhalis, D., & Sinarta, Y. (2019). Real-time co-creation and nowness service: Lessons from tourism and hospitality. *Journal of Travel & Tourism Marketing, 36*(5), 563–582. https://doi.org/10.1080/10548408.2019.1592059

Cohen, J. (1988). *Statistical power analysis for the behavioral sciences* (2nd ed.). Lawrence Erlbaum Associates. https://doi.org/10.4324/9780203771587

Davies, D. L., & Bouldin, D. W. (1979). A cluster separation measure. *IEEE Transactions on Pattern Analysis and Machine Intelligence, PAMI-1*(2), 224–227. https://doi.org/10.1109/TPAMI.1979.4766909

El Asri, L., Schulz, H., Sharma, S., Zumer, J., Harris, J., Fine, E., Mehrotra, R., & Suleman, K. (2017). Frames: A corpus for adding memory to goal-oriented dialogue systems. In *Proceedings of the 18th Annual SIGdial Meeting on Discourse and Dialogue* (pp. 207–219). Association for Computational Linguistics. https://doi.org/10.18653/v1/W17-5526

Faul, F., Erdfelder, E., Lang, A.-G., & Buchner, A. (2007). G*Power 3: A flexible statistical power analysis program for the social, behavioral, and biomedical sciences. *Behavior Research Methods, 39*(2), 175–191. https://doi.org/10.3758/BF03193146

Gilardi, F., Alizadeh, M., & Kubli, M. (2023). ChatGPT outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy of Sciences, 120*(30), Article e2305016120. https://doi.org/10.1073/pnas.2305016120

Greenhouse, S. W., & Geisser, S. (1959). On methods in the analysis of profile data. *Psychometrika, 24*(2), 95–112. https://doi.org/10.1007/BF02289823

Hennig, C. (2007). Cluster-wise assessment of cluster stability. *Computational Statistics & Data Analysis, 52*(1), 258–271. https://doi.org/10.1016/j.csda.2006.11.025

Hubert, L., & Arabie, P. (1985). Comparing partitions. *Journal of Classification, 2*(1), 193–218. https://doi.org/10.1007/BF01908075

Inaba, M., Chiba, Y., Higashinaka, R., Komatani, K., Miyao, Y., & Nagai, T. (2022). Collection and analysis of travel agency task dialogues with age-diverse speakers. In *Proceedings of the Thirteenth Language Resources and Evaluation Conference* (pp. 5759–5767). European Language Resources Association. https://aclanthology.org/2022.lrec-1.619

Ivanov, S., & Webster, C. (2019). Conceptual framework of the use of robots, artificial intelligence and service automation in travel, tourism, and hospitality companies. In *Robots, artificial intelligence, and service automation in travel, tourism and hospitality* (pp. 7–37). Emerald Publishing Limited. https://doi.org/10.1108/978-1-78756-687-320191001

Jannach, D., Manzoor, A., Cai, W., & Chen, L. (2021). A survey on conversational recommender systems. *ACM Computing Surveys, 54*(5), Article 105. https://doi.org/10.1145/3453154

Johnson, J., Douze, M., & Jégou, H. (2021). Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data, 7*(3), 535–547. https://doi.org/10.1109/TBDATA.2019.2921572

Lee, J., Chen, F., Dua, S., Cer, D., Shanbhogue, M., Naim, I., Hernández Ábrego, G., Li, Z., Chen, K., Schechter Vera, H., Ren, X., Zhang, S., Salz, D., Boratko, M., Han, J., Chen, B., Huang, S., Rao, V., Suganthan, P., . . . Duerig, T. (2025). *Gemini embedding: Generalizable embeddings from Gemini* (arXiv:2503.07891). arXiv. https://doi.org/10.48550/arXiv.2503.07891

Liu, J., Shen, D., Zhang, Y., Dolan, B., Carin, L., & Chen, W. (2022). What makes good in-context examples for GPT-3? In *Proceedings of Deep Learning Inside Out (DeeLIO 2022): The 3rd Workshop on Knowledge Extraction and Integration for Deep Learning Architectures* (pp. 100–114). Association for Computational Linguistics. https://doi.org/10.18653/v1/2022.deelio-1.10

McInnes, L., Healy, J., & Astels, S. (2017). hdbscan: Hierarchical density based clustering. *Journal of Open Source Software, 2*(11), Article 205. https://doi.org/10.21105/joss.00205

McInnes, L., Healy, J., & Melville, J. (2018). *UMAP: Uniform manifold approximation and projection for dimension reduction* (arXiv:1802.03426). arXiv. https://doi.org/10.48550/arXiv.1802.03426

Melián-González, S., Gutiérrez-Taño, D., & Bulchand-Gidumal, J. (2021). Predicting the intentions to use chatbots for travel and tourism. *Current Issues in Tourism, 24*(2), 192–210. https://doi.org/10.1080/13683500.2019.1706457

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., & Duchesnay, É. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research, 12*, 2825–2830. https://jmlr.org/papers/v12/pedregosa11a.html

Rebedea, T., Dinu, R., Sreedhar, M. N., Parisien, C., & Cohen, J. (2023). NeMo Guardrails: A toolkit for controllable and safe LLM applications with programmable rails. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing: System Demonstrations* (pp. 431–445). Association for Computational Linguistics. https://doi.org/10.18653/v1/2023.emnlp-demo.40

Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing* (pp. 3982–3992). Association for Computational Linguistics. https://doi.org/10.18653/v1/D19-1410

Rousseeuw, P. J. (1987). Silhouettes: A graphical aid to the interpretation and validation of cluster analysis. *Journal of Computational and Applied Mathematics, 20*, 53–65. https://doi.org/10.1016/0377-0427(87)90125-7

Seabold, S., & Perktold, J. (2010). Statsmodels: Econometric and statistical modeling with Python. In *Proceedings of the 9th Python in Science Conference* (pp. 92–96). SciPy. https://doi.org/10.25080/Majora-92bf1922-011

Wobbrock, J. O., Findlater, L., Gergle, D., & Higgins, J. J. (2011). The aligned rank transform for nonparametric factorial analyses using only ANOVA procedures. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems* (pp. 143–146). Association for Computing Machinery. https://doi.org/10.1145/1978942.1978963

Yao, S., Shinn, N., Razavi, P., & Narasimhan, K. (2024). *τ-bench: A benchmark for tool-agent-user interaction in real-world domains* (arXiv:2406.12045). arXiv. https://doi.org/10.48550/arXiv.2406.12045

Zang, X., Rastogi, A., Sunkara, S., Gupta, R., Zhang, J., & Chen, J. (2020). MultiWOZ 2.2: A dialogue dataset with additional annotation corrections and state tracking baselines. In *Proceedings of the 2nd Workshop on Natural Language Processing for Conversational AI* (pp. 109–117). Association for Computational Linguistics. https://doi.org/10.18653/v1/2020.nlp4convai-1.13

Zhang, T., Kishore, V., Wu, F., Weinberger, K. Q., & Artzi, Y. (2020). BERTScore: Evaluating text generation with BERT. In *Proceedings of the 8th International Conference on Learning Representations*. OpenReview. https://openreview.net/forum?id=SkeHuCVFDr

Appendix A

Repository Structure, Replication Guide, and Raw Source Schema

All code, data schemas, cached LLM outputs, and report builders are available at https://github.com/surnella/QM640-Dialogue-Clustering-Capstone. Table A1 outlines the repository directory structure, Table A2 lists the exact shell commands that reproduce every table and figure, and Table A3 provides Part A (raw source fields) and Part B (canonical schema fields) of the data dictionary.

Table A1

Repository Structure and Artifact Inventory

Directory or file path	Contents and analytical purpose
data/raw/	Downloaded Microsoft Frames (1,369 dialogues) and MultiWOZ 2.2 corpora (excluded from Git)
data/canonical/	Normalized two-tier Pydantic dialogue trajectories (canonical_corpus.jsonl)
data/interim/	Cleaned samples, partitions, embeddings, primary/refined archetypes, LLM annotations, and RQ1/RQ2 JSON summaries
data/final/	Cached Gemini 2.5 Flash/Pro JSONL responses and scored CSV/JSON results for Refined RQ2, RQ3 (N=144), and RQ4 (N=180)
src/canonical_schema.py, src/ingestion.py	Pydantic two-tier data models and corpus parsers
src/rule_engine.py, rules/rules_enterprise.yaml	Deterministic declarative policy engine and YAML rule definitions
src/final_pipeline/	Reproducible pipeline scripts for Steps 0–2i (power, cleaning, RQ1 clustering/sensitivity/refined archetypes), RefinedArchetypeRouter, Refined RQ2, RQ3 2×2 experiment, RQ4 2×3 simulation, and EDA/architecture figures
src/report/, src/build_final_report.py	APA 7 Word/PDF document generator with automated cross-reference and TOC pagination verification
notebooks/01_.. to 05_..	Five executed Jupyter notebooks covering EDA, RQ1, RQ2, RQ3, and RQ4
outputs/interim/, outputs/final/	All CSV summary tables and 300-DPI PNG figures
final/	Submission-ready QM_640_Final_Report_Sandeep_Surnella.docx and .pdf

Table A2

End-to-End Replication Commands

Pipeline stage	Command (run from repository root with .venv active)	Outputs produced
0. Power analysis	python src/final_pipeline/step0_sample_size.py	Table 8
1. Cleaning & univariate EDA	python src/final_pipeline/step1_prepare_data.py	Tables 5, 6, E1; Figures 1, E1
2. Primary RQ1 clustering	python src/final_pipeline/step2_rq1_clustering.py	Table 12; Tables D1–D4; Figures D1–D3
2b–2i. RQ1 sensitivity & refined	python src/final_pipeline/step2g_rq1_truncation.py && python src/final_pipeline/step2h_rq1_ablation.py && python src/final_pipeline/step2i_rq1_refined_archetypes.py	Table 13; Tables C1–C2, D5–D6, E5; Figures D4, E3
3. Final harness & routing test	python src/final_pipeline/test_step1_harness.py	Table 13 routing note (t_rq1_hybrid_routing_diagnostics.csv)
4. Refined RQ2 experiment	python src/final_pipeline/step2_rq2_refined.py --execute	Table 14 (refined rows); Table E3; Figure 6
5. RQ3 2×2 factorial experiment	python src/final_pipeline/step3_rq3_experiment.py --execute	Table 15; Table C3; Figure 7
6. RQ4 2×3 turn-economy sim.	python src/final_pipeline/step4_rq4_simulation.py --execute	Table 16; Table C3; Figure 8
7. Bivariate/multivariate EDA & arch.	python src/final_pipeline/step5_figures_and_eda.py	Table 7; Table E2; Figures 2, 3, 4, 5
8. Build Final Report (.docx/.pdf)	python src/build_final_report.py	final/QM_640_Final_Report_Sandeep_Surnella.docx & .pdf

Note. All random seeds are fixed at 42. Cached JSONL responses allow offline execution without API calls.

Table A3

Data Dictionary, Parts A and B: Raw Corpus Fields and Canonical Schema Mapping

Source / Layer	Field name	Data type (Unit)	Description and mapping role
Frames (Raw)	id, user_id, wizard_id	String (ID)	Unique dialogue ID and role-player IDs (11 users, 12 wizards)
Frames (Raw)	turns[].author, text, timestamp	String / Float (ms)	Speaker (user or wizard), raw utterance text, and Unix timestamp
Frames (Raw)	turns[].labels.acts[].name, args	List / Key–Value	Dialogue act tags (inform, request, offer, switch_frame) and slot–value pairs
Frames (Raw)	labels.wizardSurveyTaskSuccessful, userSurveyRating	Boolean / Float (1–5)	Wizard task completion flag and user satisfaction rating

Source / Layer	Field name	Data type (Unit)	Description and mapping role
MultiWOZ (Raw)	dialogue_id, services	String / List	Unique dialogue ID and active service domains (hotel, train, taxi, attraction)
MultiWOZ (Raw)	turns[].speaker, utterance, frames[].state	String / Dict	Speaker (USER or SYSTEM), utterance text, and turn-level slot state dictionary
Canonical (Tier 1)	origin_city, destination_city, accommodation_tier	String (City / Tier)	Normalized origin, destination, and hotel star rating or category
Canonical (Tier 1)	departure_date, return_date, duration_nights	String (Date) / Int (nights)	Normalized trip dates and length of stay in nights
Canonical (Tier 1)	traveler_count_adults, traveler_count_children, budget_amount	Int (persons) / Float (USD)	Party composition and numeric budget ceiling in US dollars
Canonical (Tier 2)	visa_processing_window_days, passport_expiry_months	Int (days / months)	Departure lead window (days) and remaining passport validity (months)
Canonical (Tier 2)	package_type, sharing_basis, discount_applied_percentage	String / Float (%)	Package class (FIT or Group), room occupancy, and promotional discount (%)

Appendix B

Ten-Week Project Schedule and Milestone Completion Log

Table B1 documents the completed 10-week execution schedule aligned with QM640 course milestones.

Table B1

Ten-Week Project Schedule and Milestone Completion Status

Week	QM640 course topic	Completed project tasks and deliverables	Status
1	Define research problem & topic	Business & technical problem statements finalized; GitHub repository initialized; Frames & MultiWOZ 2.2 ingested	Complete
2	Synopsis draft	Four research questions, hypotheses, and two-tier Pydantic schema defined	Complete
3	Literature search & data collection	26 peer-reviewed business & technical sources synthesized; canonical corpus parsed (N = 2,834 analytic)	Complete
4	Objectives & hypotheses	Data quality audit, univariate/bivariate/multivariate EDA, noncentral-F power analysis, and primary RQ1 clustering	Complete
5	Research design & methodology	Interim report submitted; primary RQ2 experiment (99 calls) and RQ1 sensitivity/truncation grid (Steps 2b–2i) completed	Complete
6	Data analysis plan	Built RefinedArchetypeRouter (Blocks prefix router, κ = .535–.958); repeated RQ2 on refined archetypes (p < .001)	Complete
7	Data analysis plan execution	Executed RQ3 2 × 2 factorial experiment (N = 144) and dual-LLM judge reliability audit (gemini-2.5-pro, κ = .917)	Complete
8	Data analysis assessment	Executed RQ4 2 × 3 multi-turn simulation (N = 180), persona fidelity validation (95.6%, κ = .933), and HC3/ART checks	Complete
9	Refinement & documentation	Generated publication architecture, workflow, and EDA diagrams; synchronized empirical Jupyter notebooks and README	Complete
10	Document results & findings	Compiled, audited, and exported submission-ready Final Report (.docx and .pdf)	Complete

Appendix C

LLM Behavioral Annotation Codebook, Reliability, and Dual-Judge Validation

Table C1 defines all 16 structured fields of the LLM behavioral codebook (`UserBehaviour` in *step2e_llm_user_annotation.py*: 10 ordinal 1–5 rating scales, 3 categorical goal/tone clustering fields, and 3 audit-only fields excluded from clustering), Table C2 reports test–retest reliability across all 16 fields and criterion validity against ground-truth Frames dialogue acts, and Table C3 summarizes the cross-model dual-judge reliability in RQ3 and simulator persona validation in RQ4.

Table C1

LLM Behavioral Annotation Codebook (All 16 Structured Schema Fields)

# & Schema field	Scale & role	Operational definition and scale anchors
1. destination_flexibility	Ordinal 1–5 (Cluster)	1 = insists on one fixed destination; 5 = open to any destination recommendation
2. date_flexibility	Ordinal 1–5 (Cluster)	1 = fixed travel dates; 5 = wide date window or completely flexible timing
3. exploration_breadth	Ordinal 1–5 (Cluster)	1 = evaluates a single option; 5 = actively compares multiple destinations, hotels, or dates
4. decisiveness	Ordinal 1–5 (Cluster)	1 = hesitant or leaves undecided; 5 = commits quickly and firmly once a matching offer is shown
5. detail_orientation	Ordinal 1–5 (Cluster)	1 = asks no package details; 5 = probes amenities, star ratings, flight times, and inclusions
6. revision_tendency	Ordinal 1–5 (Cluster)	1 = never modifies initial constraints; 5 = repeatedly shifts dates, destinations, or budgets
7. question_asking	Ordinal 1–5 (Cluster)	1 = passive responder; 5 = asks frequent clarifying and comparison questions
8. politeness	Ordinal 1–5 (Cluster)	1 = curt or abrupt; 5 = consistently courteous and appreciative
9. price_sensitivity	Ordinal 1–5 (Cluster)	1 = price never mentioned or irrelevant; 5 = price/budget dominates every decision
10. information_upfront	Ordinal 1–5 (Cluster)	1 = reveals constraints gradually across turns; 5 = states nearly all trip constraints in opening message
11. primary_goal	Categorical (Cluster)	6 classes: explore_destinations, price_shopping, specific_trip_booking, compare_options, information_only, other
12. secondary_goals	Multi-label (Cluster)	Multi-hot subset of the 6 goal classes above clearly present in user utterances

# & Schema field	Scale & role	Operational definition and scale anchors
13. emotional_tone	Categorical (Cluster)	4 classes: positive, neutral, impatient, or frustrated
14. outcome	Categorical (Audit)	4 classes: booked, no_suitable_option, undecided_or_left, unclear (excluded from clustering input)
15. party_type	Categorical (Audit)	6 classes: solo, couple, family, group, business, unknown (excluded from clustering input)
16. unusual_conversation	Boolean (Audit)	True only if off-task or anomalous request (excluded from clustering input)

Note. Fields 1–13 form the clustering feature matrix (*llm_matrix*: 10 ordinal + 6 primary_goal + 6 secondary_goals + 4 emotional_tone = 26 numeric columns). Fields 14–16 are retained strictly for post-clustering validation.

Table C2

Test–Retest Reliability and Criterion Validity Across All 16 LLM Behavioral Schema Fields

Schema field (#1–16)	Test–retest κ	Ground-truth Frames act criterion or field role	Spearman ρ ($p < .001$)
1. Destination flexibility	.96	—	—
2. Date flexibility	.89	—	—
3. Exploration breadth	.95	Frame switches per user turn	.30
		Comparison requests per user turn	.26
4. Decisiveness	.91	—	—
5. Detail orientation	.89	Distinct preference slots mentioned	.45
6. Revision tendency	.89	Slot-value changes in user acts	.73
7. Question asking	.90	Request acts per user turn	.65
8. Politeness	.91	—	—
9. Price sensitivity	.95	Numeric budget stated	.28
10. Information given up front	.87	Share of slots stated in first user turn	.32
11. Primary goal (6 classes)	.78	Categorical Cohen's κ (clustering)	—
12. Secondary goals (multi-label)	.75	Binary presence Cohen's κ (clustering)	—
13. Emotional tone (4 classes)	.88	Categorical Cohen's κ (clustering)	—
14. Session outcome (4 classes)	.98	Categorical Cohen's κ (audit-only)	—
15. Party type (6 classes)	.93	Categorical Cohen's κ (audit-only)	—
16. Unusual conversation (boolean)	.80	Boolean Cohen's κ (99.0% raw agreement, audit-only)	—

Note. Test–retest subsample $n = 100$ (quadratic-weighted κ for ordinal Fields 1–10; Cohen's κ for Fields 11–16); criterion validity $N = 1{,}334$ Frames dialogues.

Table C3

Cross-Model Dual-Judge Reliability (RQ3) and Simulator Persona Fidelity Validation (RQ4)

Validation audit	Sample *n*	Evaluator models compared	Raw agreement / Accuracy	Cohen's κ (PABAK)	Additional validity evidence
RQ3 Policy Compliance Dual-Judge Audit	48	*gemini-2.5-flash* vs. *gemini-2.5-pro*	95.8%	.917 (.917)	Flash vs. deterministic regex audit: 89.6% agreement, κ = .791
RQ4 Simulator Persona Fidelity Audit	45	Intended Archetype vs. Blind *gemini-2.5-pro* Judge	95.6% (vs. 33.3% chance)	.933 (p < .001)	Mean ratings (Decisiveness / Exploration / Rigidity): DB = 5.0/1.0/1.5; OE = 3.9/4.3/2.0; RCS = 4.5/1.1/4.3

Appendix D

RQ1 Sensitivity Grid, Ablation Diagnostics, and Primary Archetype Profiles

Tables D1–D6 and Figures D1–D4 provide the complete diagnostic tables and visual plots supporting the RQ1 primary clustering benchmark and the representation × truncation sensitivity analysis.

Table D1

Choice of Number of Clusters k on the Golden Partition Across Initial MiniLM and Final 46-d Blocks (M = 9) Representations

Representation & *k*	Cosine Silhouette	Davies–Bouldin index	Inertia / Bootstrap ARI	Smallest cluster (%) & diagnostic
Panel A (MiniLM-L6-v2 whole): *k* = 2	.206	3.136	Inertia = 193.1	35.5%
Panel A (MiniLM-L6-v2 whole): *k* = 3	.141	3.244	Inertia = 184.6	24.1%
Panel A (MiniLM-L6-v2 whole): *k* = 4	.142	2.997	Inertia = 178.1	17.7%
Panel A (MiniLM-L6-v2 whole): *k* = 5	.148	2.950	Inertia = 173.6	8.7%
Panel A (MiniLM-L6-v2 whole): *k* = 6	.077	3.118	Inertia = 169.7	8.4%
Panel B (46-d Blocks, *M* = 9): *k* = 2	.112 (cal. .044)	2.796	ARI = .949 (96.2% core)	40.4% — Collapses Decisive & Rigid into 1 cluster
Panel B (46-d Blocks, *M* = 9): *k* = 3	.094 (cal. .033)	2.647	ARI = .888 (93.1% core)	25.5% — Adopted winner (lowest DB, ARI >= .60, min >= 20%)
Panel B (46-d Blocks, *M* = 9): *k* = 4	.066 (cal. .015)	2.756	ARI = .589 (55.8% core)	12.5% — Fails stability (ARI < .60) & min size (< 20%)
Panel B (46-d Blocks, *M* = 9): *k* = 5	.069 (cal. .026)	2.784	ARI = .606 (61.1% core)	11.4% — Fails stability (ARI < .60) & min size (< 20%)
Panel B (46-d Blocks, *M* = 9): *k* = 6	.065 (cal. .019)	2.729	ARI = .631 (53.2% core)	11.3% — Fragmented min cluster (11.3% < 20% floor)

Note. On the final 46-d *Blocks* (*M* = 9) representation (Panel B), *k* = 3 achieves the lowest (best) Davies–Bouldin index (2.647), high bootstrap stability (ARI = .888 in the truncation grid; .869 in

the 6-config ablation), and balanced cluster sizes (≥ 25.5%), whereas $k = 2$ collapses *Decisive Booker* and *Rigid-Constraint Seeker* and $k \geq 4$ fails stability/size thresholds.

Table D2

Pairwise Post Hoc Comparisons of Mean Silhouette Across Clustering Algorithms

Algorithm comparison	Mean difference	95% CI (Tukey HSD)	*p* (Tukey HSD)	*p* (Games–Howell)
K-Means − Agglomerative	.029	[.013, .046]	< .001	< .001
HDBSCAN − Agglomerative	.093	[.076, .110]	< .001	< .001
HDBSCAN − K-Means	.063	[.047, .080]	< .001	< .001

Table D3

Behavioral Profiles of the Three Primary Whole-Dialogue Archetypes (Golden Partition)

Feature	Direct Transactional	Constraint Negotiator	Guided Explorer
n (%)	240 (25.7%)	225 (24.1%)	468 (50.2%)
Turns, *M* (*SD*)	9.92 (5.02)	15.64 (6.84)	17.19 (6.10)
Words per user / agent turn	10.60 / 11.67	12.94 / 18.15	9.63 / 15.71
Slot-value revisions / Hesitations	6.9 / 0.7	20.3 / 1.1	16.7 / 1.1
Agent package offers	0.84	4.28	3.73
Budget stated (%) / Task success (%)	63.7% / 92.1%	70.7% / 94.2%	75.4% / 95.3%

Table D4

Primary Archetype Assignment and Centroid Similarity Across Partitions and MultiWOZ 2.2

Sample	*N*	Centroid cosine *M* (*SD*)	Direct Transactional (%)	Constraint Negotiator (%)	Guided Explorer (%)
Frames golden	933	.895 (.042)	25.7	24.1	50.2
Frames held-out	401	.894 (.042)	29.4	20.2	50.4
MultiWOZ 2.2	1,499	.805 (.037)	0.9	15.1	84.0

Table D5

Stability, Routing Agreement, and Outcome Association of Hybrid Archetypes by Conversation Length

User messages used (*M*)	Dialogues shortened (%)	Bootstrap ARI	Core members (%)	Smallest cluster (%)	Raw L1 κ	Blocks router κ	Outcome *V*
First 3 messages	94.7	.89	92.7	18.3	.72	.359	.16
First 5 messages	78.2	.84	77.1	20.3	.38	.535	.27
First 7 messages	53.2	.85	89.9	29.7	.25	.722	.38
First 9 messages	30.7	.89	93.1	25.5	.24	.958	.45
Whole conversation	0.0	.62	53.7	7.9	.22	—	.57

Note. Raw L1 κ uses unprojected 5-turn features (step2g); Blocks router κ projects prefix features into the exact golden Blocks space (RefinedArchetypeRouter).

Table D6

Ablation of the Refined Archetype Definition at M = 9 User Messages (k = 3)

Representation configuration	ARI (whole)	ARI (9 msgs)	Core members (%)	Outcome *V*	ARI vs. adopted
Gemini embedding + Gemini behavior profile (adopted)	.56	.87	89.2	.45	1.00
MiniLM embedding + Gemini behavior profile	.69	.85	87.7	.40	.33
Gemini behavior profile only	.73	.85	97.6	.35	.17
Gemini embedding only	.91	.91	92.0	.34	.45
MiniLM embedding only	.68	.72	72.1	.27	.14
MiniLM embedding + dialogue-act features (no LLM)	.61	.76	88.0	.45	.16

Figure D1

Cluster-Number Diagnostics on Golden Partition

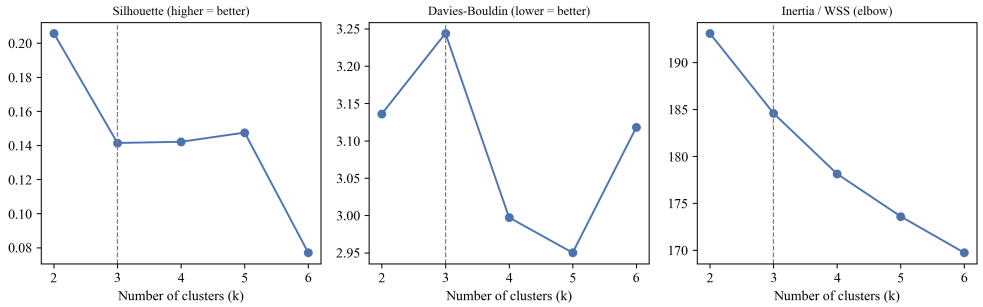


Figure D2

Silhouette Coefficients Across 159 Subsample Runs by Algorithm

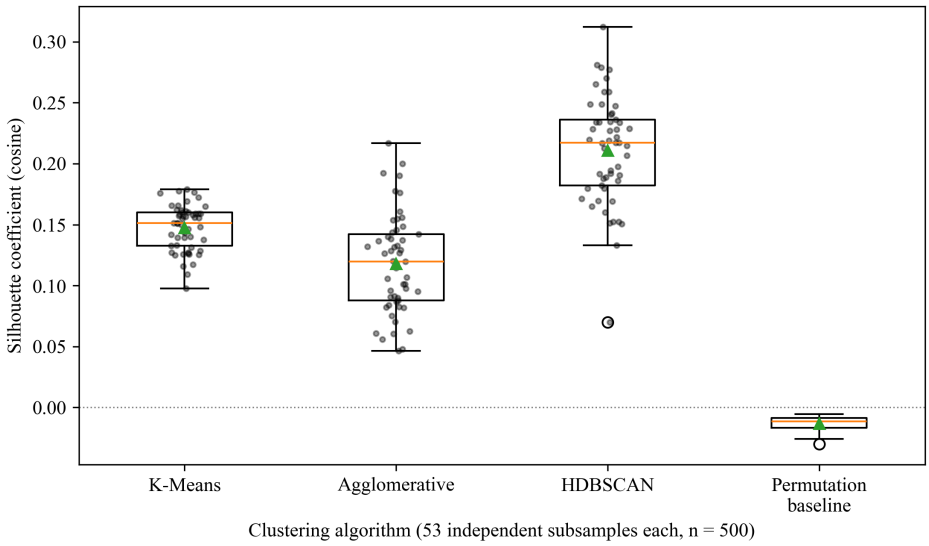


Figure D3

UMAP Projection of Frames Archetypes and MultiWOZ 2.2 External Transfer Check

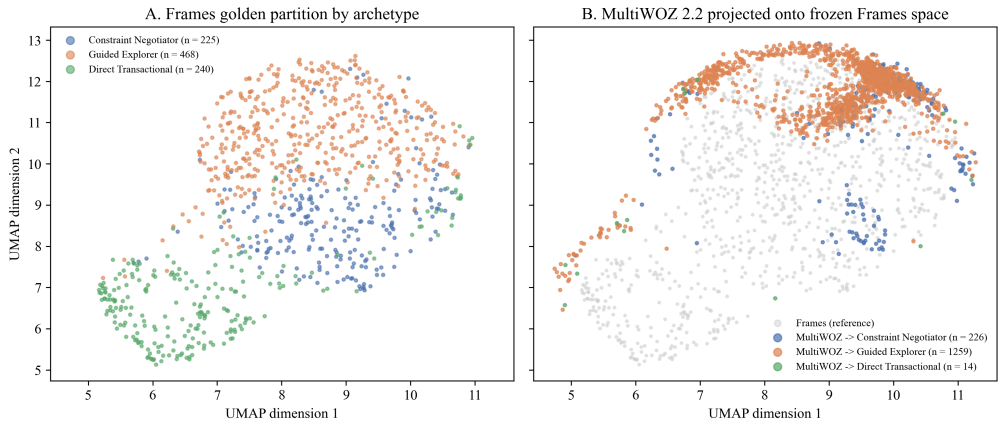
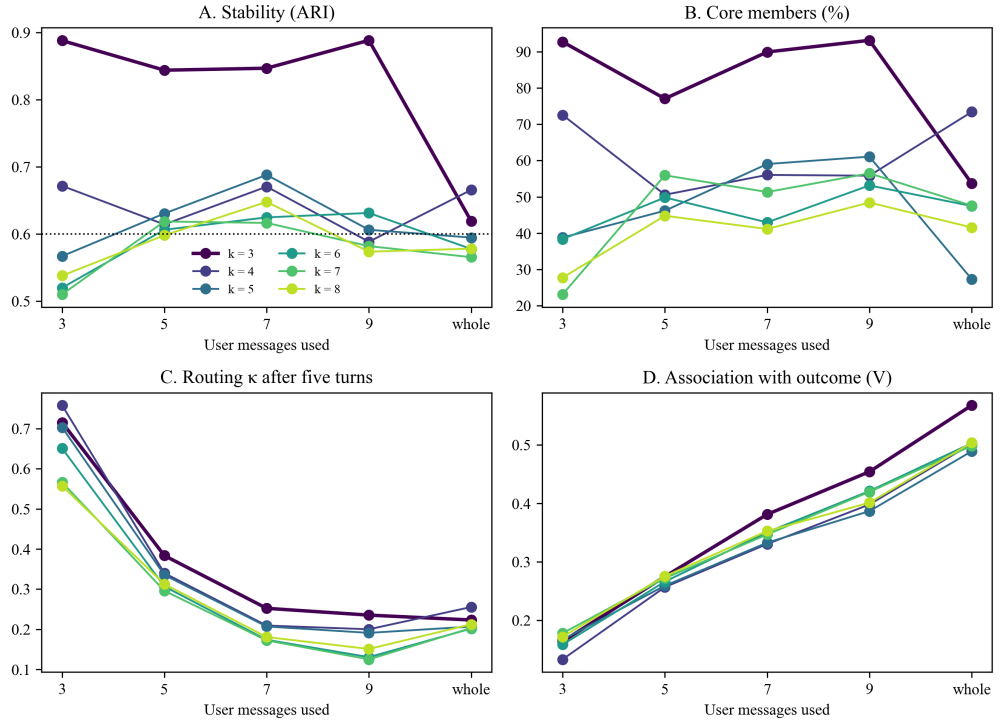


Figure D4

Hybrid Archetype Stability, Routing, and Outcome Association Across Truncation Horizons



Appendix E

Supplementary EDA, RQ2 Pairwise Contrasts, and RQ3/RQ4 Groundwork Diagnostics

Tables E1–E5 and Figures E1–E3 provide the supplementary slot-completeness table, multivariate PCA loadings, Holm-adjusted pairwise contrasts for RQ2, corpus rule coverage audit for RQ3, and human turn-economy baseline for RQ4.

Table E1

Observed Canonical Slot Completeness by Corpus

Canonical slot	Schema tier	Microsoft Frames (%)	MultiWOZ 2.2 (%)
origin city	Tier 1	99.4	70.8
destination city	Tier 1	99.0	71.4
departure date	Tier 1	84.9	68.6
return date	Tier 1	70.8	0.0
duration nights	Tier 1	61.8	27.0
traveler count adults	Tier 1	63.9	54.0
traveler count children	Tier 1	18.4	0.0
budget amount	Tier 1	72.9	24.1
accommodation tier	Tier 1	72.7	19.7
visa processing window days	Tier 2	0.0	0.0
passport expiry months	Tier 2	0.0	0.0
discount applied percentage	Tier 2	0.0	0.0

Table E2

Multivariate Principal Component Loadings of Structural and Behavioral Features (Frames)

Variable	PC1 loading (40.8% var.)	PC2 loading (17.2% var.)	PC3 loading (10.6% var.)
Total Turns	.393	.166	−.234
Slot Revisions	.424	.195	−.064
Agent Offers	.361	.292	−.078
User Words/Turn	.025	−.087	.888
Dest. Flexibility	.327	−.182	.135
Exploration Breadth	.417	−.108	.101
Decisiveness	−.259	.503	.028
Revision Tendency	.297	−.460	.025
Question Asking	.301	.248	.027
Politeness	.089	.520	.341

Table E3

Holm-Adjusted Paired Contrasts for Refined RQ2 Outcomes (N = 33)

Outcome metric	Pairwise contrast	Mean diff.	*t*	*df*	*p* (Holm)	d_z
Refined: SBERT style cosine	Cluster-Adaptive vs. Static	.144	4.24	32	< .001	0.74
Refined: SBERT style cosine	Cluster-Adaptive vs. Zero-Shot	.067	2.36	32	.049	0.41
Refined: SBERT style cosine	Static vs. Zero-Shot	−.077	−2.20	32	.049	−0.38
Refined: Gemini style cosine	Cluster-Adaptive vs. Zero-Shot	.027	1.20	32	.713	0.21
Refined: Gemini style cosine	Static vs. Zero-Shot	.024	0.79	32	.865	0.14
Refined: Gemini style cosine	Cluster-Adaptive vs. Static	.004	0.13	32	.895	0.02
Refined: BERTScore F1	Cluster-Adaptive vs. Static	.056	1.45	32	.472	0.25
Refined: BERTScore F1	Cluster-Adaptive vs. Zero-Shot	.025	0.85	32	.804	0.15
Refined: BERTScore F1	Static vs. Zero-Shot	−.031	−0.75	32	.804	−0.13
Refined: Pacing error (words)	Cluster-Adaptive vs. Zero-Shot	−6.25	−3.49	32	.004	−0.61
Refined: Pacing error (words)	Cluster-Adaptive vs. Static	−2.46	−2.66	32	.024	−0.46
Refined: Pacing error (words)	Static vs. Zero-Shot	−3.79	−1.82	32	.078	−0.32
Refined: Response length (words)	Cluster-Adaptive vs. Static	6.82	7.00	32	< .001	1.22
Refined: Response length (words)	Static vs. Zero-Shot	−10.91	−4.29	32	< .001	−0.75
Refined: Response length (words)	Cluster-Adaptive vs. Zero-Shot	−4.09	−1.84	32	.075	−0.32

Table E4

Empirical Corpus Audit of Policy Rule Coverage in Microsoft Frames (N = 1,334)

Declarative rule	Required canonical slot	Slot observed (%)	Dialogues triggering (*n*)	Triggering (%)
pax capacity	traveler count adults	24.1	54	4.05
budget minimum	budget amount	65.1	266	19.94
visa window	visa processing window days	0.0	0	0.00
passport validity	passport expiry months	0.0	0	0.00
discount threshold	discount applied percentage	0.0	0	0.00

Table E5

Human Wizard Baseline of Turn Economy and Redundant Elicitation by Refined Archetype (Frames)

Refined user archetype	*n*	Whole-dialogue turns *M* (*SD*)	RER *M* (*SD*)	Wizard request turns *M*
Decisive Booker	595	14.56 (4.52)	.014 (.047)	1.44
Open Explorer	337	20.75 (7.07)	.016 (.049)	1.82
Rigid-Constraint Seeker	402	10.60 (5.59)	.017 (.061)	1.27

Note. Turns in Table E5 reflect total observed human-wizard turns (user + wizard messages) across complete archival Frames transcripts (including post-offer discussion and non-booking endings), whereas Table 16 (RQ4) measures total conversational turns until package agreement (<code>[ACCEPT_BOOKING]</code>) in closed-loop simulation.

Figure E1

Observed Slot Completeness by Corpus

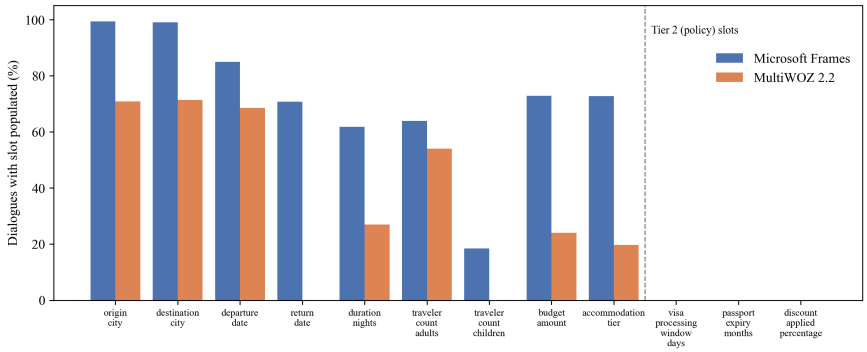


Figure E2

Primary RQ2 Results on Unrefined Whole-Dialogue Archetypes

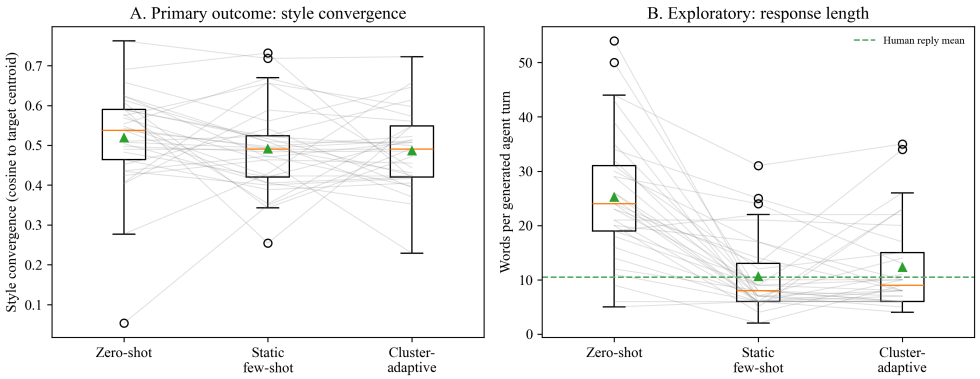
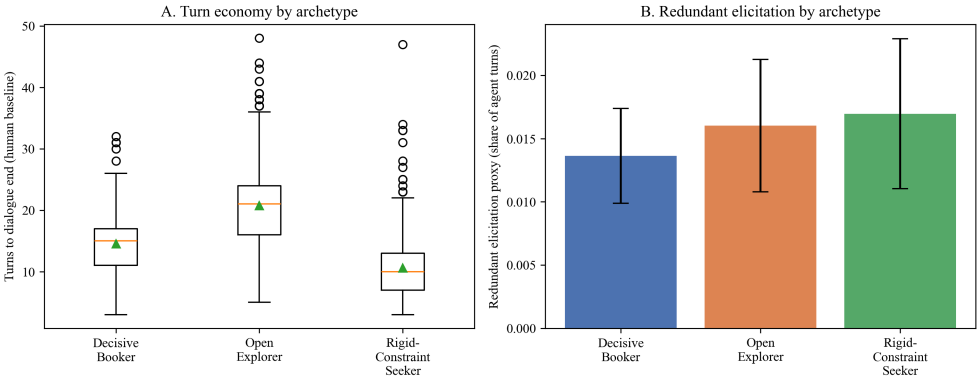


Figure E3

Human Wizard Baseline of Turn Economy and Redundant Elicitation by Refined Archetype



Appendix F

Derivation of User Archetype Cluster Names and Deterministic Centroid-Mapping Protocol

Unsupervised clustering algorithms such as *K*-Means output anonymous integer cluster identifiers (<code>0</code>, <code>1</code>, <code>2</code>) rather than semantic persona names, and those integer identifiers are invariant up to permutation across random initializations (the *label-switching problem*). To ensure that human-interpretable archetype labels are assigned reproducibly without manual post-hoc cherry-picking, both the primary whole-dialogue clustering pipeline (*src/final_pipeline/step2_rq1_clustering.py*) and the refined opening-prefix clustering pipeline (*src/final_pipeline/step2i_rq1_refined_archetypes.py*) apply pre-declared, deterministic Python centroid-mapping rules to the golden-partition cluster means. Table F1 documents the exact sequential decision rule, distinguishing centroid criterion, and resulting golden-partition profile for each primary and refined archetype.

Table F1

Deterministic Centroid-Mapping Rules for Primary and Refined User Archetypes

Clustering taxonomy & script	Step & deterministic Python rule on cluster centroid means	Assigned archetype name	Distinguishing golden-partition centroid signature ($N = 933$)
Primary Whole-Dialogue Archetypes (*step2_rq1_clustering.py*)	Step 1: <code>idxmax(turns_M)</code> among the 3 clusters	Guided Explorer	Longest dialogues ($M = 19.65$ turns, $n = 187$, 20.0%); highest option exploration across whole transcripts
Primary Whole-Dialogue Archetypes (*step2_rq1_clustering.py*)	Step 2: <code>idxmax(revisions)</code> among the remaining 2 clusters	Constraint Negotiator	Highest slot-revision count ($M = 1.98$ revisions, $M = 15.80$ turns, $n = 351$, 37.6%)
Primary Whole-Dialogue Archetypes (*step2_rq1_clustering.py*)	Step 3: Remaining cluster ID	Direct Transactional	Shortest whole-dialogue sessions ($M = 11.98$ turns, $n = 395$, 42.3%); lowest revision count

Clustering taxonomy & script	Step & deterministic Python rule on cluster centroid means	Assigned archetype name	Distinguishing golden-partition centroid signature ($N = 933$)
Refined Opening-Prefix Archetypes (*step2i_rq1_refined_archetypes.py*)	Step 1: <code>idxmax(destination_flexibility)</code> across the 3 *Blocks* clusters ($M = 9$)	Open Explorer	Highest destination flexibility (3.49/5.0) and exploration breadth (2.96/5.0); longest dialogues (20.78 turns; $n = 238$, 25.5%)
Refined Opening-Prefix Archetypes (*step2i_rq1_refined_archetypes.py*)	Step 2: <code>idxmax(decisiveness)</code> among the remaining 2 clusters	Decisive Booker	Highest decisiveness (3.97/5.0), politeness (4.26/5.0), and specific trip booking goal (99.3%; $n = 420$, 45.0%)
Refined Opening-Prefix Archetypes (*step2i_rq1_refined_archetypes.py*)	Step 3: Remaining cluster ID	Rigid-Constraint Seeker	Lowest destination/date flexibility (2.20/2.18) and detail orientation (1.46/5.0); highest non-booking rate under rigid constraints (68.0% no suitable option; $n = 275$, 29.5%)

Note. All centroid-mapping rules are evaluated strictly on golden-partition means (D_{gold}, $n = 933$) after fitting K-Means ($k = 3$, seed = 42). Held-out test dialogues (D_{test}, $n = 401$) inherit the mapped archetype name of their nearest golden centroid.

Appendix G

Role, Architecture, and Validation Scope of Judge and Annotator LLMs

A critical methodological principle of this study is that **all primary hypothesis tests (H_{01} through H_{04}) are anchored in deterministic mathematical and programmatic Python verification metrics** rather than subjective LLM preference scores. Specifically, RQ1 cluster quality is evaluated via mathematical Silhouette, Davies–Bouldin, Calinski–Harabasz, and Adjusted Rand Index formulas; RQ2 style convergence and pacing calibration are evaluated via deterministic non-Gemini encoders (*all-MiniLM-L6-v2* cosine, *roberta-large* BERTScore F_1, and exact word-count error); RQ3 exemplar entity leakage (ELR) and slot grounding F_1 are computed by deterministic Python string/set verifiers (*src/final_pipeline/metrics.py*); and RQ4 turn economy (turns to agreed package quotation and Redundant Elicitation Rate) is computed by deterministic turn counters and regex slot-ledger matching.

Within this architecture, LLMs serve two strictly bounded measurement and validation functions: (a) in RQ1, *gemini-2.5-flash* (at temperature = 0) acts as a **structured behavioral codebook annotator** prior to *K*-Means clustering, extracting the 16-field *UserBehaviour* schema (validated via test–retest κ_w = .87–.96 and ground-truth Frames dialogue-act correlations ρ = .30–.73; Table C2); and (b) in RQ3 and RQ4, the larger *gemini-2.5-pro* model serves strictly as an **independent cross-verification and quality-control validator** (Table C3). Importantly, the independent Judge LLM does not invent or contribute any standalone performance metric to firm up the paper's statistical conclusions; rather, it acts as an methodological gatekeeper. During the RQ3 pilot phase, cross-checking outputs against the deterministic verifier identified a prompt-ordering sensitivity (*pilot_v1* → *pilot_v2*) that was resolved prior to full execution. In the final experiments, because the independent *gemini-2.5-pro* validator strongly corroborated both the

RQ3 policy evaluations (95.8% agreement, κ = .917 with *gemini-2.5-flash*; 93.1% agreement, κ = .861 with the deterministic regex verifier) and the RQ4 simulated customer personas (95.6% blind classification accuracy, κ = .933), no manual re-annotation or rework was required. Table G1 summarizes the exact role, deterministic counterpart, and gatekeeping outcome of every LLM annotator and judge in the pipeline.

Table G1

Functional Scope, Primary Deterministic Metrics, and Cross-Verification Role of Judge and Annotator LLMs Across RQ1–RQ4

Research question & pipeline stage	Primary hypothesis-test metrics (Deterministic / Mathematical)	LLM annotator / Judge model & role	Cross-verification agreement & gatekeeping outcome
RQ1: Opening-Prefix Behavioral Annotation (*step2e_llm_user_annotation.py*)	Cosine Silhouette, Davies–Bouldin, Calinski–Harabasz, Bootstrap ARI (.89 truncation grid / .87 ablation), and *Blocks* prefix-to-cluster κ (.359–.958)	**Codebook Annotator (*gemini-2.5-flash*, temp = 0):** Extracts Fields 1–13 (clustering inputs) and Fields 14–16 (audit-only) before *K*-Means runs; does not name clusters.	Test–retest κ_w = .87–.96 (mean .91) on ordinal Fields 1–10; κ = .75–.98 on Fields 11–16; significant Spearman ρ = .30–.73 ($p < .001$) against ground-truth Frames acts (Table C2).
RQ2: Style Convergence & Pacing Calibration (*step2_rq2_refined.py*)	Hybrid & pure SBERT style cosine (*all-MiniLM-L6-v2*), BERTScore F_1 (*roberta-large*), and absolute word-count error ($\lvert Words_{gen} - Words_{human} \rvert$)	**No Judge LLM used:** All evaluation encoders are deterministic, non-Gemini neural and lexical metrics.	Eliminates same-family LLM evaluation bias in RQ2; RM-ANOVA and Friedman tests confirm $p < .001$ on both hybrid and pure SBERT style convergence.
RQ3: Policy Compliance Dual-Judge Audit (*step3_rq3_experiment.py*)	Deterministic regex rule verifier (*det_pass*), exact Exemplar Entity Leakage Rate (ELR), and deterministic Slot Grounding F_1	**Primary Judge (*gemini-2.5-flash*) + Independent Validator Judge (*gemini-2.5-pro*):** Cross-verifies natural-language enforcement of 5 enterprise policy rules on stratified audit subsample (n = 48).	Guided *pilot_v1* → *pilot_v2* prompt ordering fix; final dual-judge raw agreement = 95.8%, Cohen's κ = .917 (PABAK = .917), and 93.1% agreement (κ = .861) with deterministic regex verification (Table C3).

Research question & pipeline stage	Primary hypothesis-test metrics (Deterministic / Mathematical)	LLM annotator / Judge model & role	Cross-verification agreement & gatekeeping outcome
RQ4: Blind Simulator Persona Fidelity Audit (*step4_rq4_simulation.py*)	Deterministic turn counter to agreed package quotation (*total_turns*, *log_turns*), regex Redundant Elicitation Rate (RER), and goal completion rate	**Independent Blind Validator (*gemini-2.5-pro*):** Reads simulated customer transcripts (n = 45, 15 per archetype) blinded to true labels to verify persona fidelity; contributes zero metrics to RQ4 ANOVA.	Classified 43/45 simulated users into their intended archetype (95.6% accuracy vs. 33.3% chance, Cohen's κ = .933, p < .001; Table C3), validating the multi-turn simulation environment.

Note. All LLM annotation and judge calls use deterministic decoding (temperature = 0.0, Pydantic JSON schema enforcement) and are cached in *data/interim/* and *data/final/* for 100% offline reproducibility.